

Memory 2: Paging, Caching, and TLBs

Sam Kumar

CS 162: Operating Systems and System Programming


Lecture 16

<https://inst.eecs.berkeley.edu/~cs162/su20>

Read: A&D Ch 9.1-4

Recall: Address Space

- Definition: **Set of accessible addresses and the state associated with them**
 - $2^{32} = \sim 4$ billion on a 32-bit machine
- What happens when you read or write to an address?
 - Perhaps acts like regular memory
 - Perhaps causes I/O operation
 - (Memory-mapped I/O)
 - Causes program to abort (segfault)?
 - Communicate with another program
 - ...



Recall: Important Aspects of Memory Multiplexing


- Protection
 - Prevent access to private memory of other process or kernel
- Translation
 - Gives uniform view of memory to programs
 - Allows for efficient “tricks”
 - E.g., in implementation of `fork()`
- Controlled Overlap
 - Read-only data, execute-only shared libraries
 - Inter-process communication

Recall: Interposing on Process Behavior

- OS interposes on process' I/O operations
 - How? All I/O happens via syscalls.
- OS interposes on process' CPU usage
 - How? Interrupt lets OS preempt current thread
- **Question: How can the OS interpose on process' memory accesses?**
 - Too slow for the OS to interpose *every* memory access
 - Translation: hardware support to accelerate the common case
 - Page fault: uncommon cases trap to the OS to handle



Recall: Uniprogramming

- No Translation or Protection
 - Application always runs at same place in physical memory since only one application at a time
 - Application can access any physical address
 - Application given illusion of dedicated machine by giving it reality of a dedicated machine




Recall: Primitive Multiprogramming

- Multiprogramming without Translation or Protection
- Use Loader/Linker: Adjust addresses while program loaded into memory (loads, stores, jumps)
 - Everything adjusted to memory location where OS put program
 - Translation done by a linker-loader (relocation)
- **No protection!**




Recall: Multiprogramming with Protection

- Can we protect programs from each other without translation?
 - **Yes: Base and Bound!**
 - **Used by, e.g., Cray-1 supercomputer**

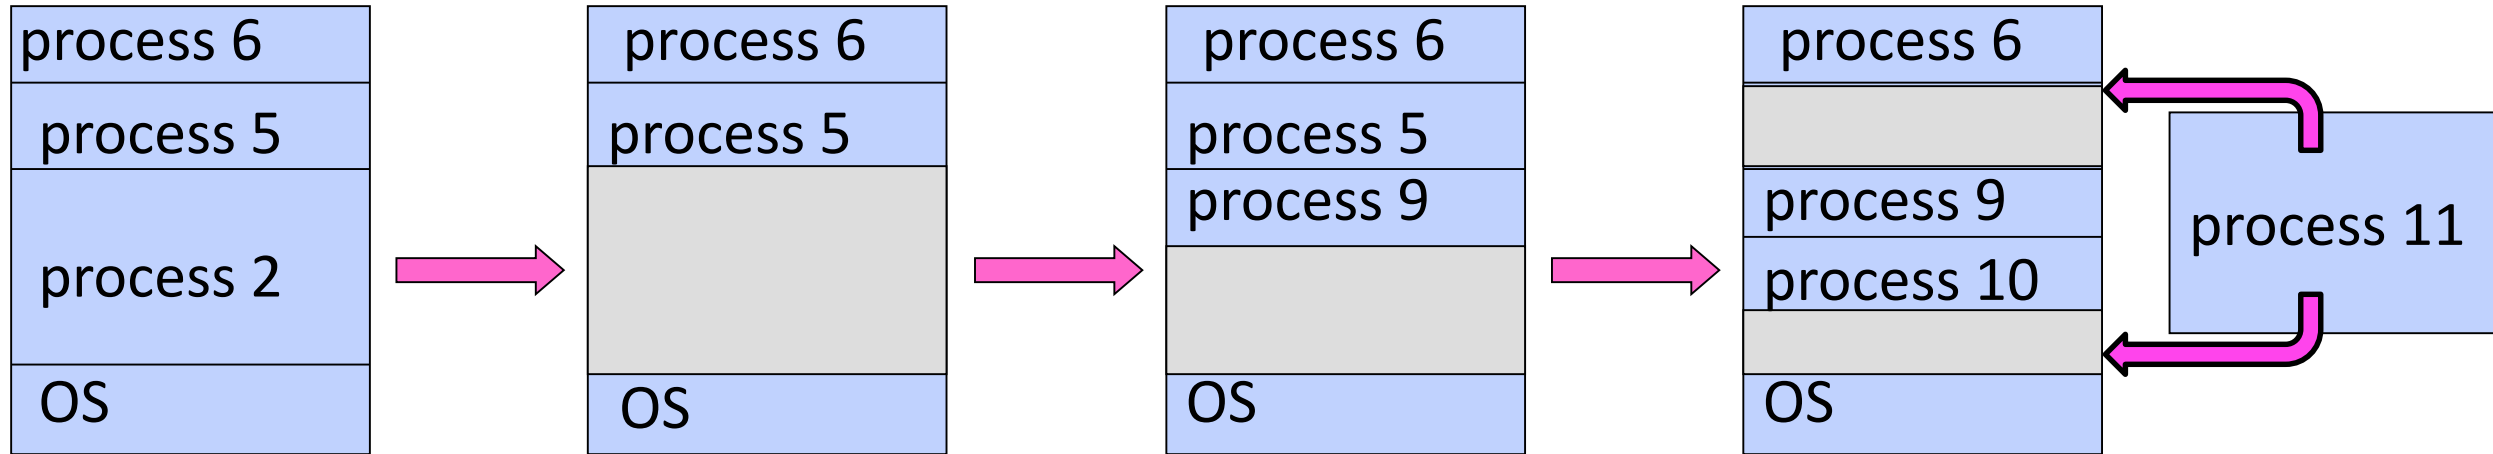


Recall: General Translation





- Two views of memory:
 - View from the CPU (what program sees, virtual memory)
 - View from memory (physical memory)
 - Hardware translator (Memory Management Unit or MMU) converts between the two views
- With translation, every program can be linked/loaded into same region of *user* address space

Recall: Issues with Simple Base and Bound



- Fragmentation problem over time
- No support for sparse address space
- Hard to do interprocess sharing
 - E.g., to share code

Recall: Segmentation

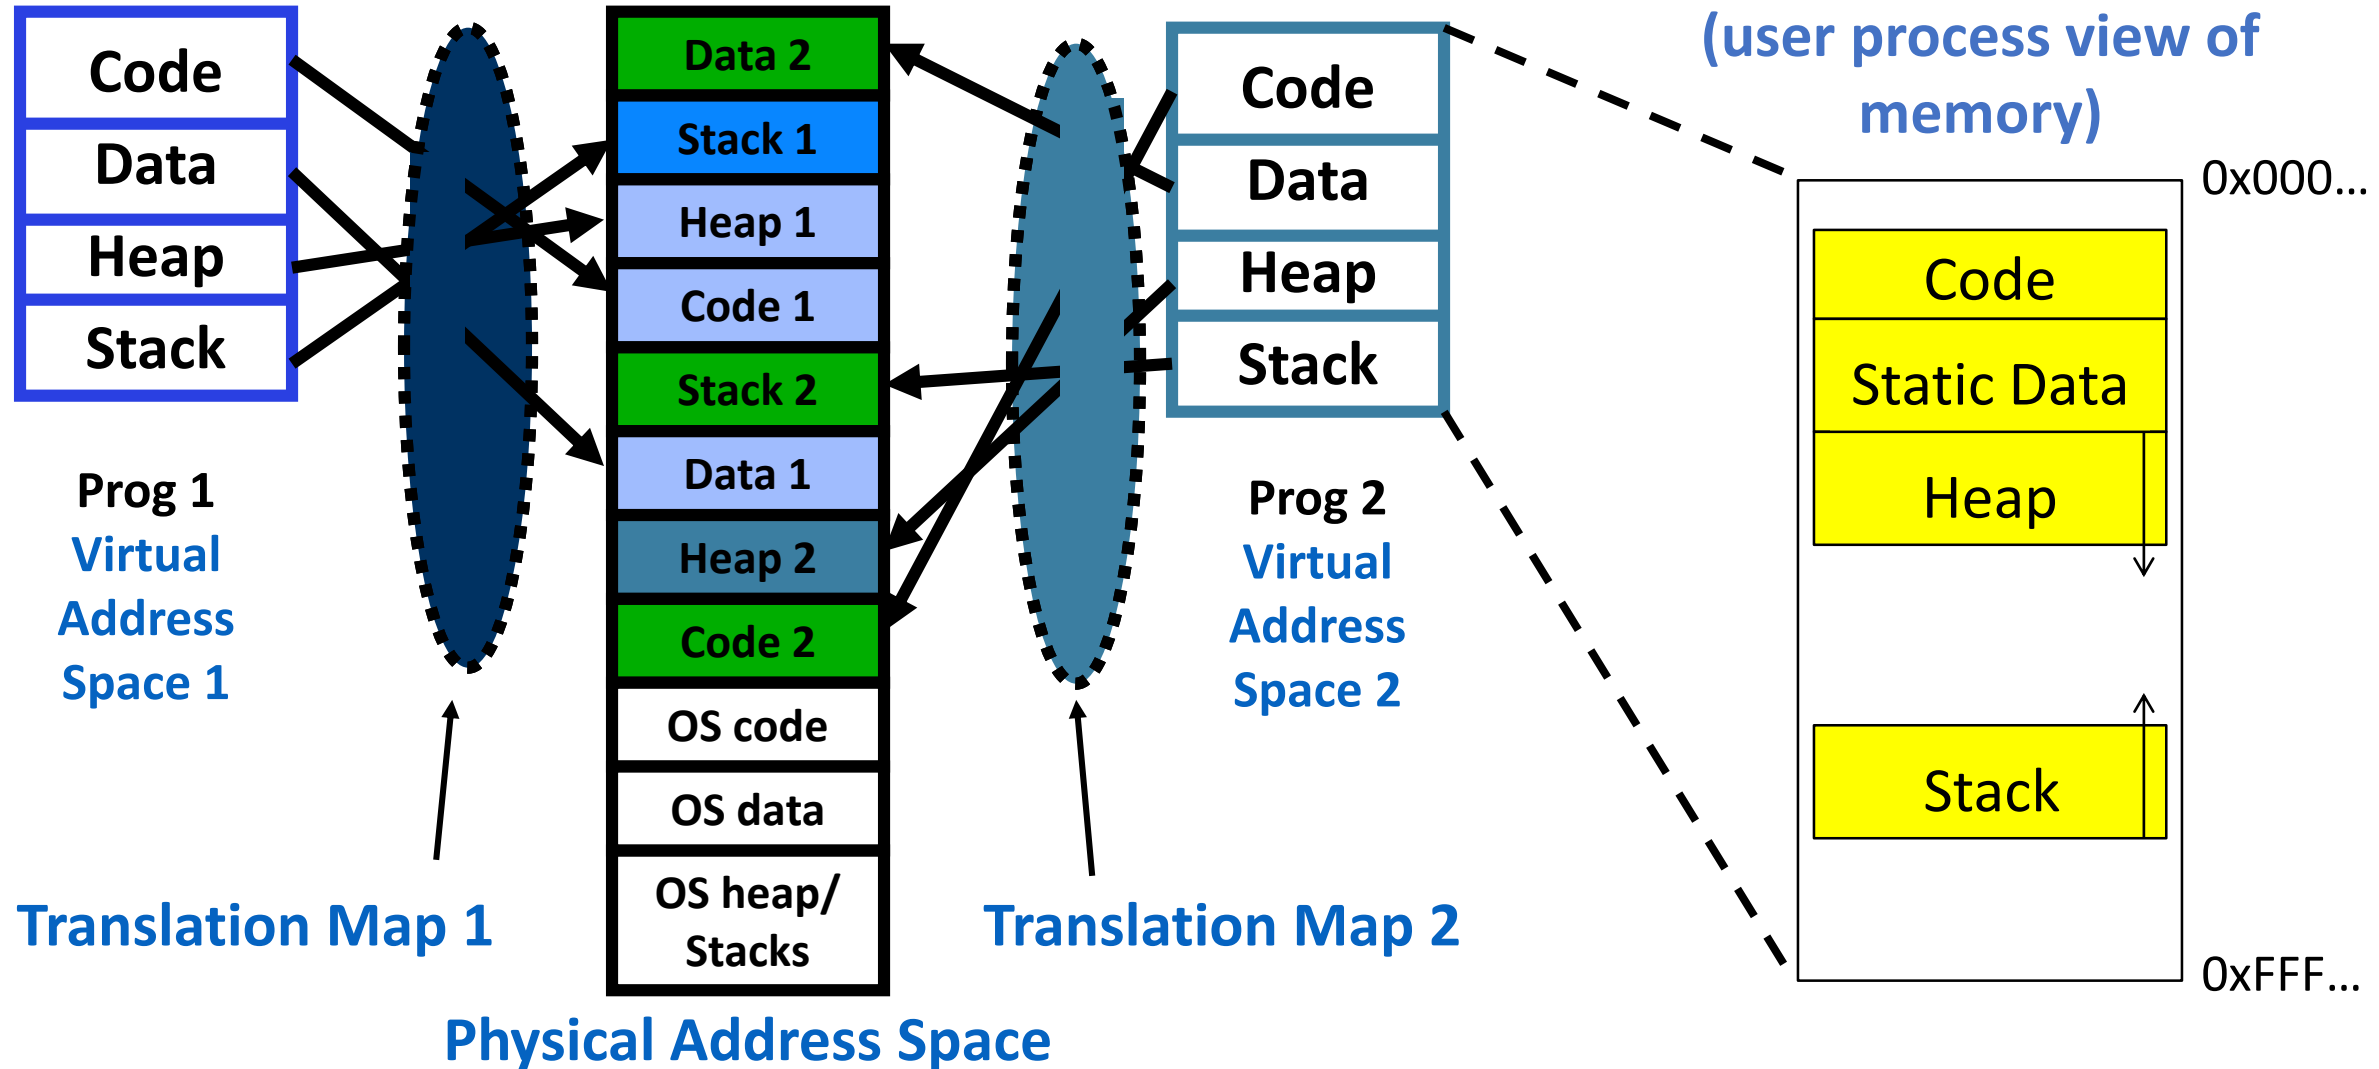


- Program's view of memory: multiple separate segments
- Each segment is given a region of contiguous memory
 - Has a base and limit
- Memory address consists of segment ID and offset

Recall: Problems with Segmentation

- Must fit variable-sized chunks into physical memory
- May move processes multiple times to fit everything
- Limited options for swapping to disk
- **Fragmentation**: wasted space
 - **External**: free gaps between allocated chunks
 - **Internal**: don't need all memory within allocated chunks

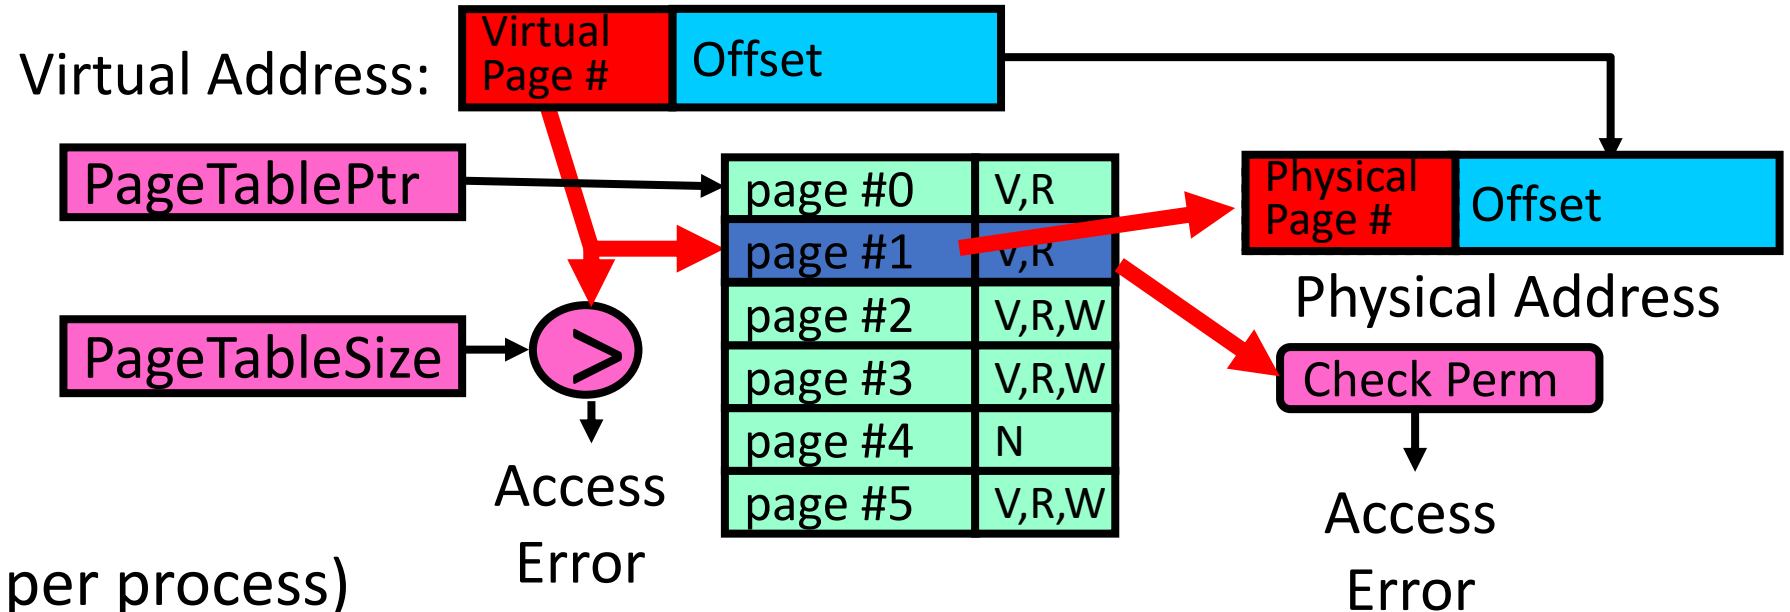
Recall: General Address Translation



Paging: *Fixed-Size* Chunks of Memory

- Divide up physical memory into equal-size chunks called *page frames*
- Divide up virtual memory into equal-size chunks called *pages*
- Key idea: each physical page frame can contain any page
- No external fragmentation!
- Should pages be as big as our previous segments?
 - No: Can lead to lots of internal fragmentation
 - Typically have small pages (1K-16K)
 - Consequently: need multiple pages/segment

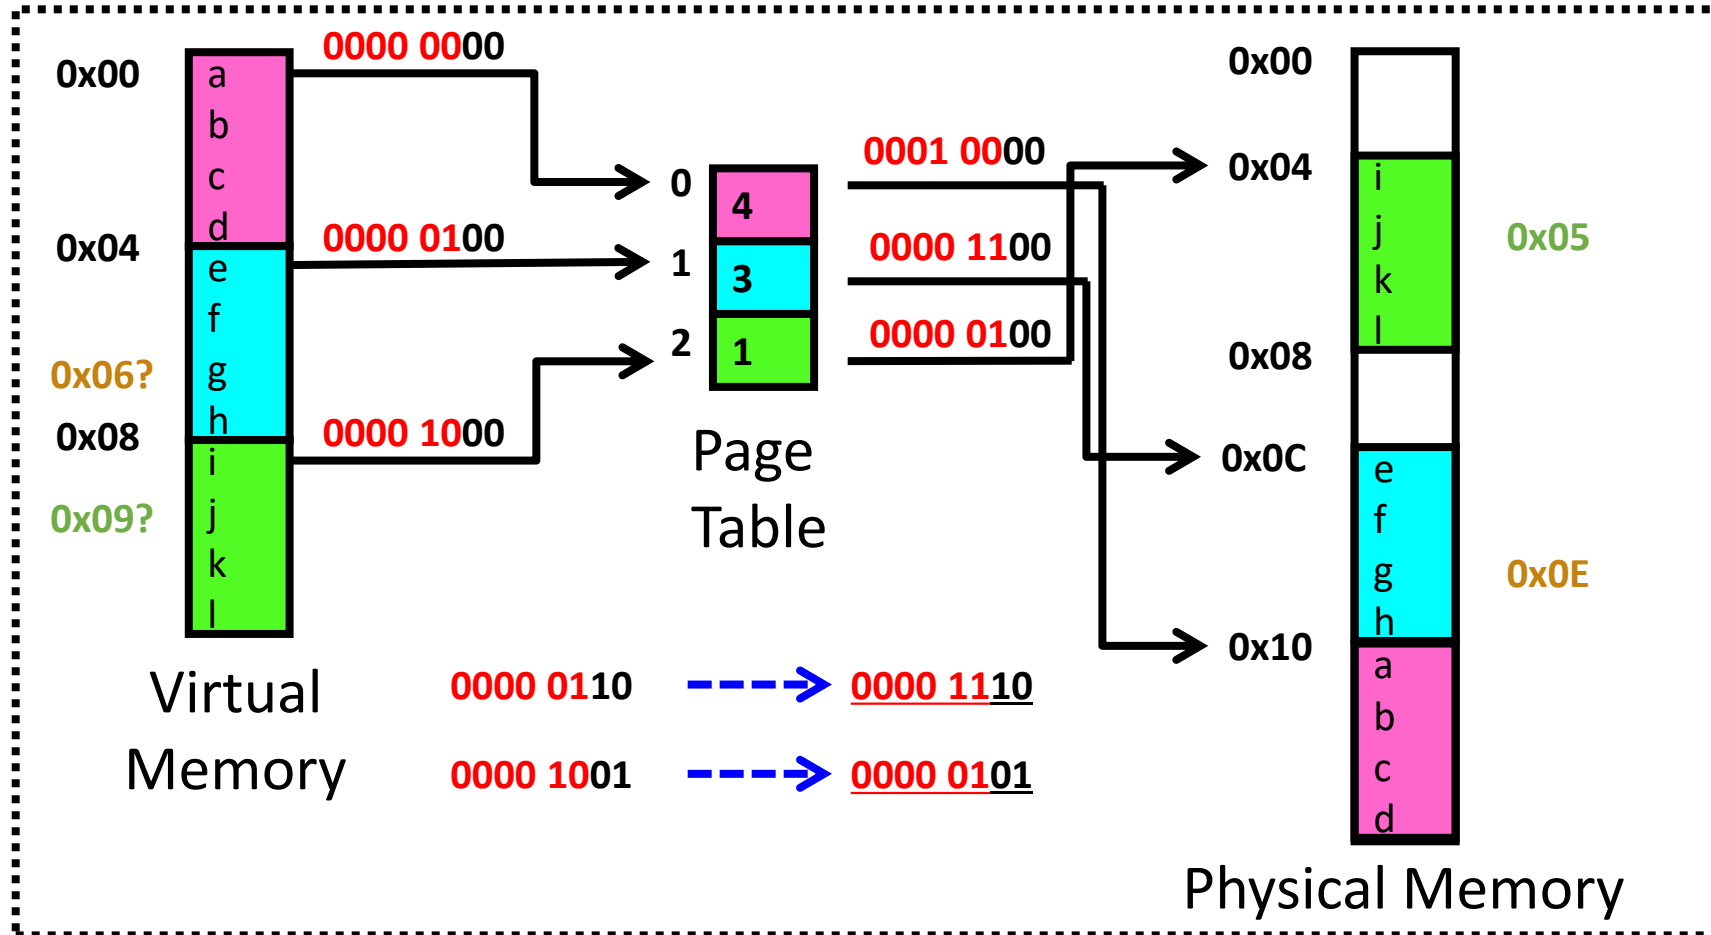
Hardware Support for Paging



- Page Table (One per process)
 - Contains physical page and permission for each virtual page
 - Permissions include: Valid bits, Read, Write, etc
- Virtual address mapping
 - Offset from Virtual address copied to Physical Address
 - Virtual page # is the *index* into the page table
 - Physical page # copied from table into physical address

Simple Page Table Example

Example (4 byte pages)



- What is the physical address for...
 - Virtual address 0x6?
 - Virtual address 0x9?

What About Sharing?

Virtual Address
(Process A):




PageTablePtrA

page #0	V,R
page #1	V,R
page #2	V,R,W
page #3	V,R,W
page #4	N
page #5	V,R,W

PageTablePtrB

page #0	V,R
page #1	N
page #2	V,R,W
page #3	N
page #4	V,R
page #5	V,R,W

Virtual Address
(Process B):

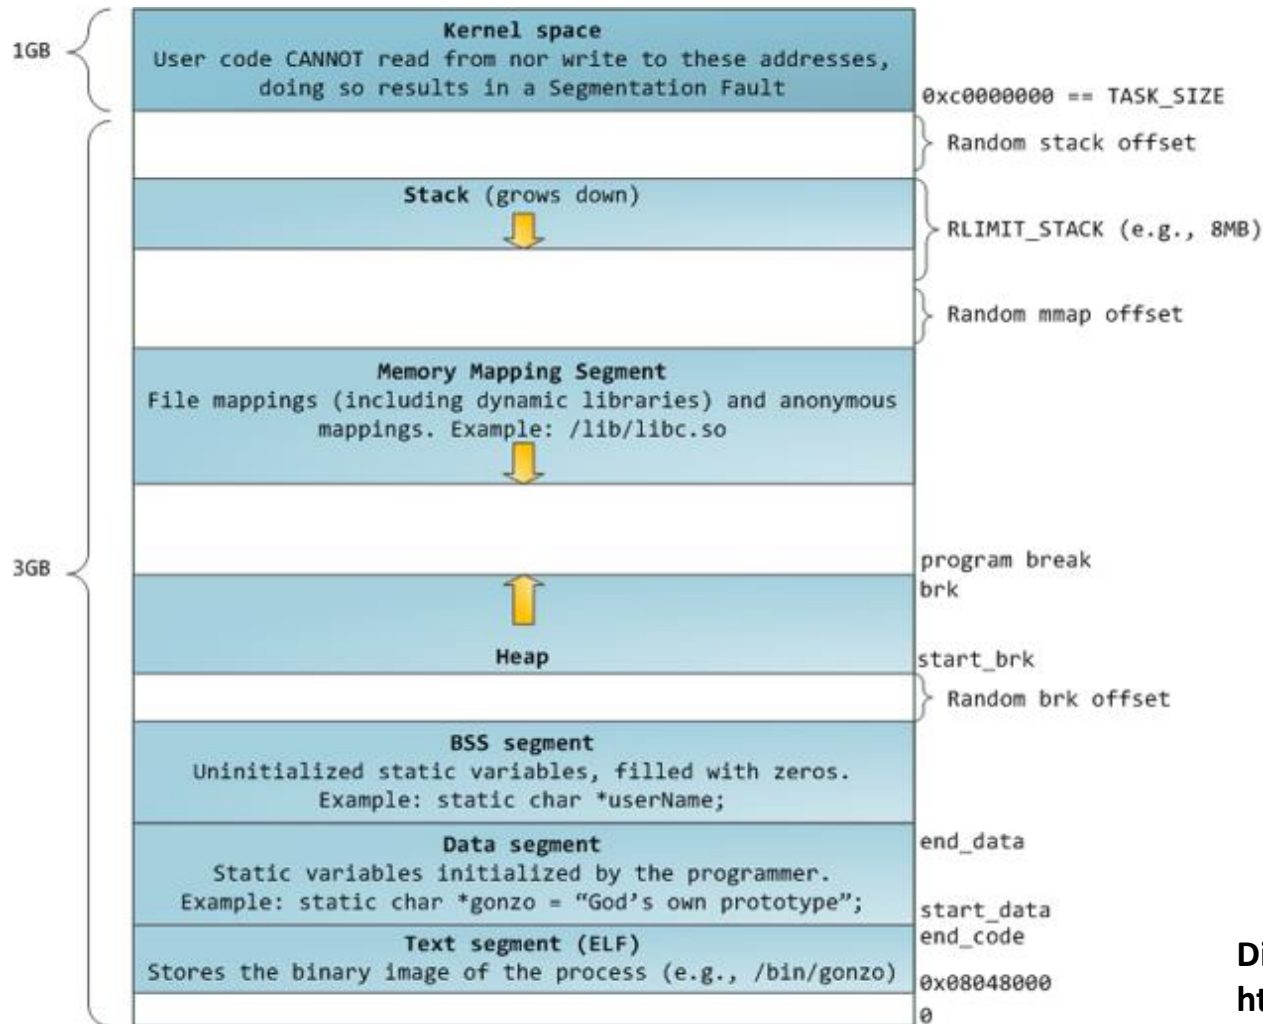


This physical page appears in both process' address spaces

Where is page sharing used?

- Kernel data mapped into each process
- Different processes running the same binary
- User-level system libraries
- Shared pages as IPC

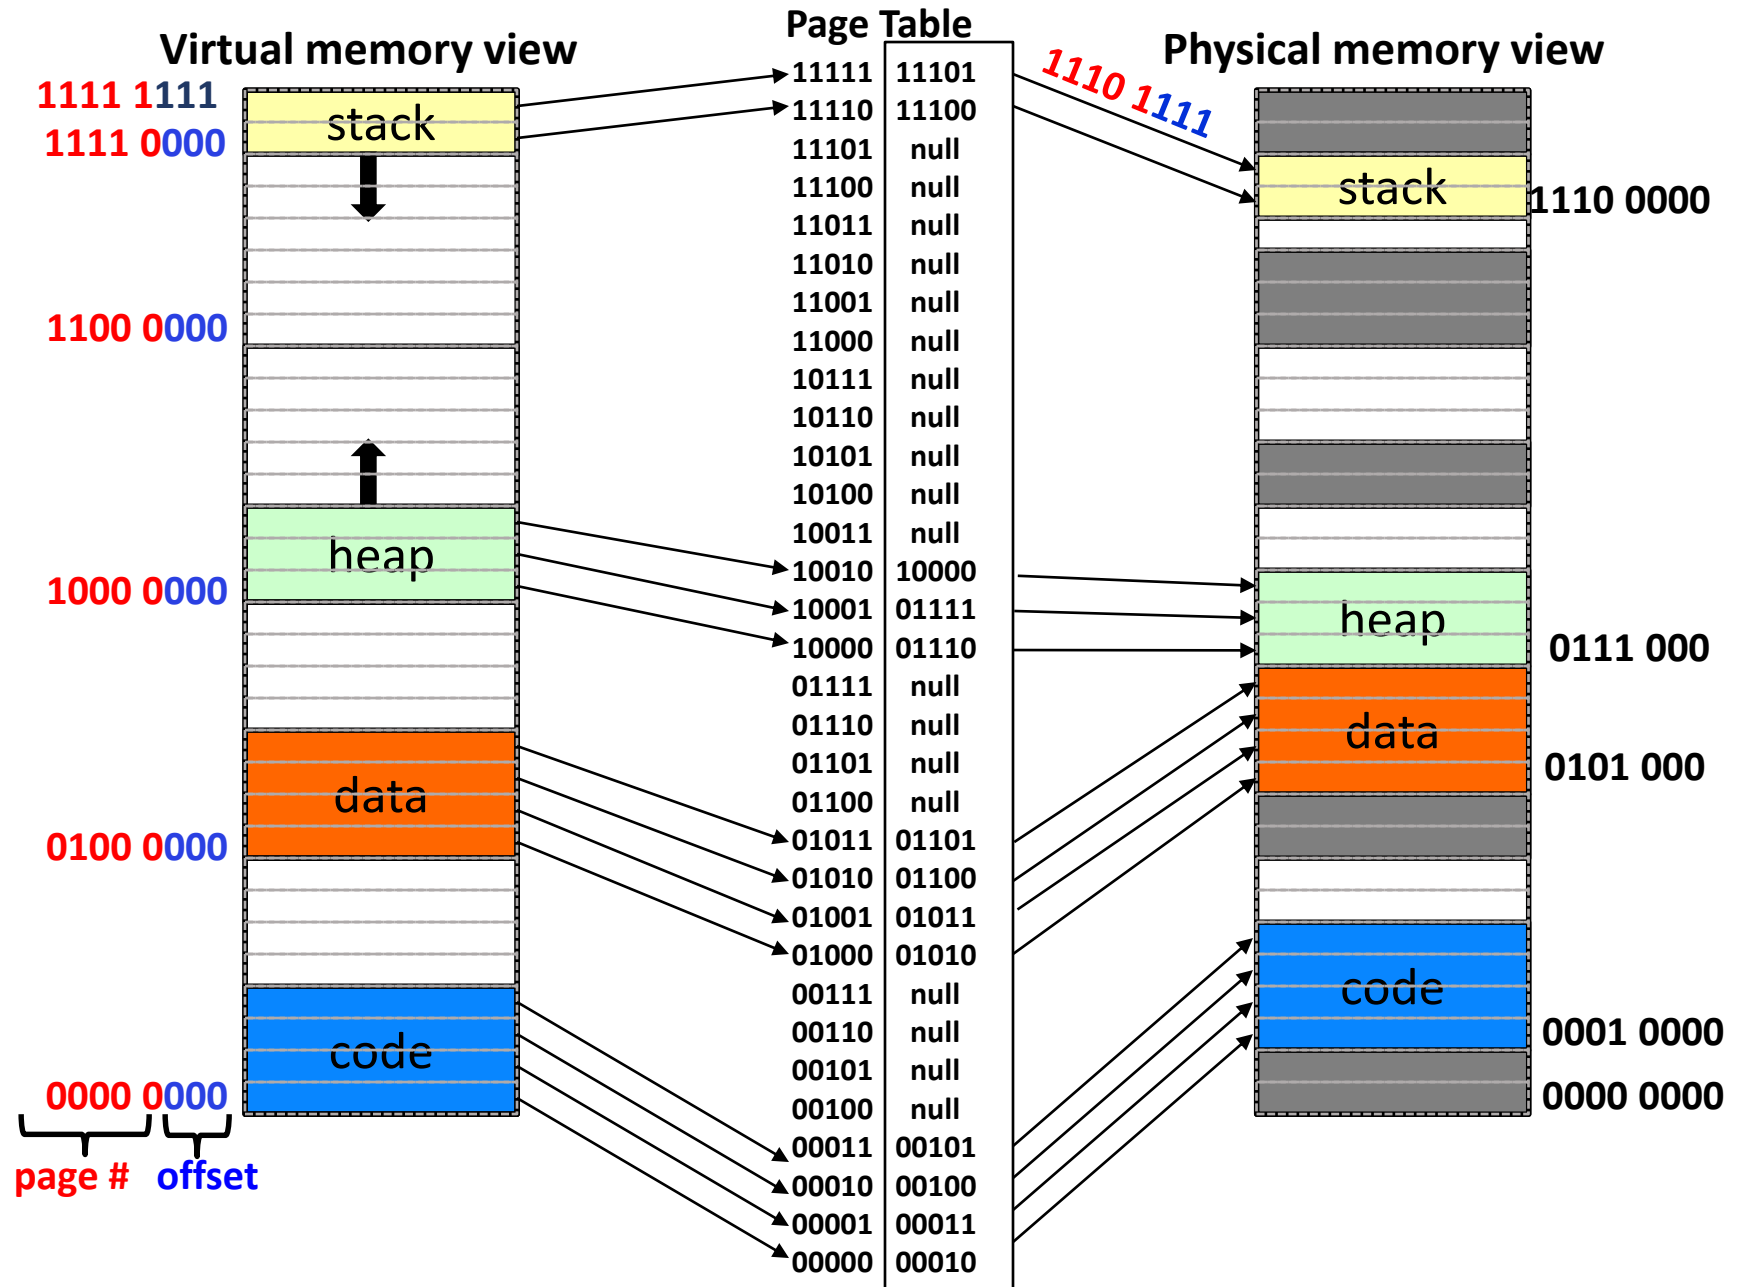
Example: Memory Layout for Linux 32-bit*



* Pre-Meltdown patches, more later...

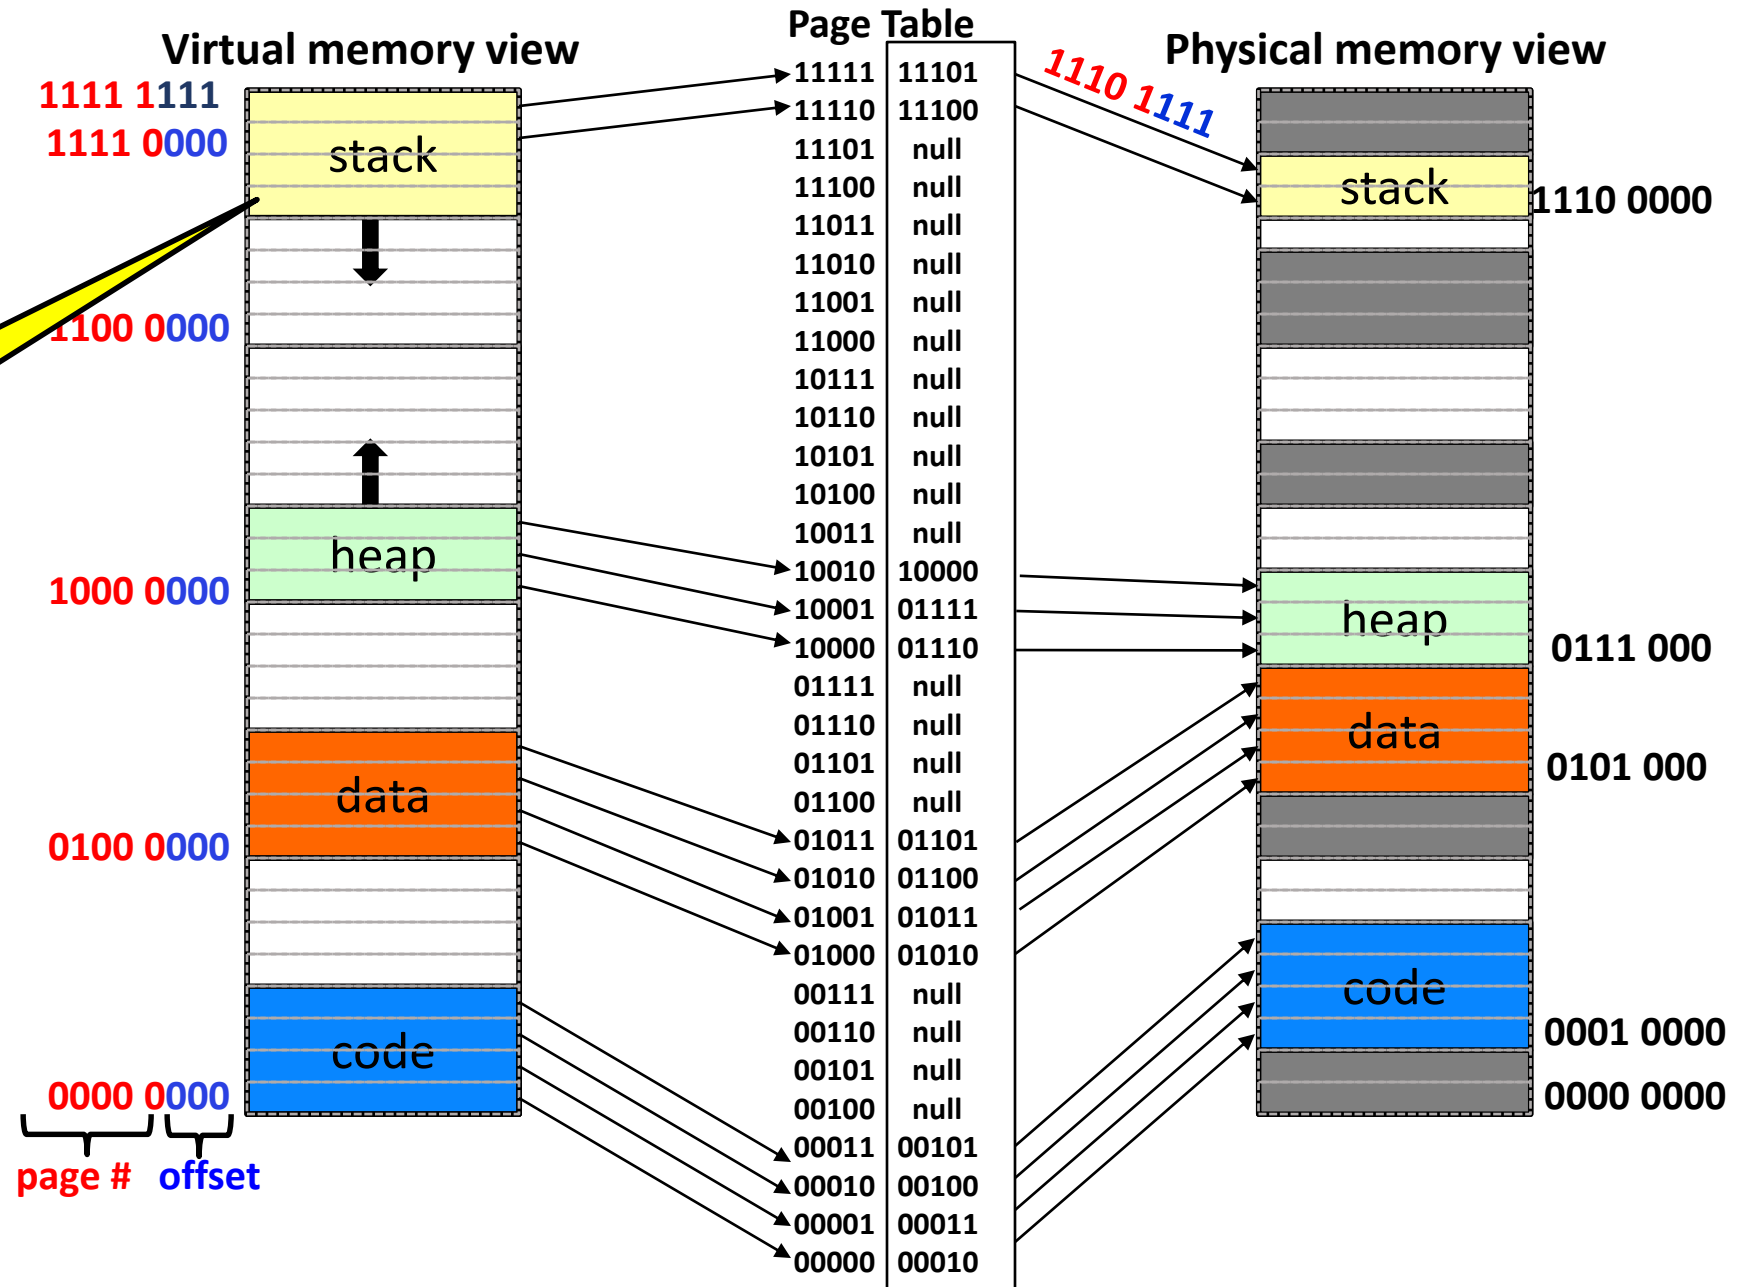
Diagram source:
<http://static.duartes.org/img/blogPosts/linuxFlexibleAddressSpaceLayout.png>

Summary: Paging



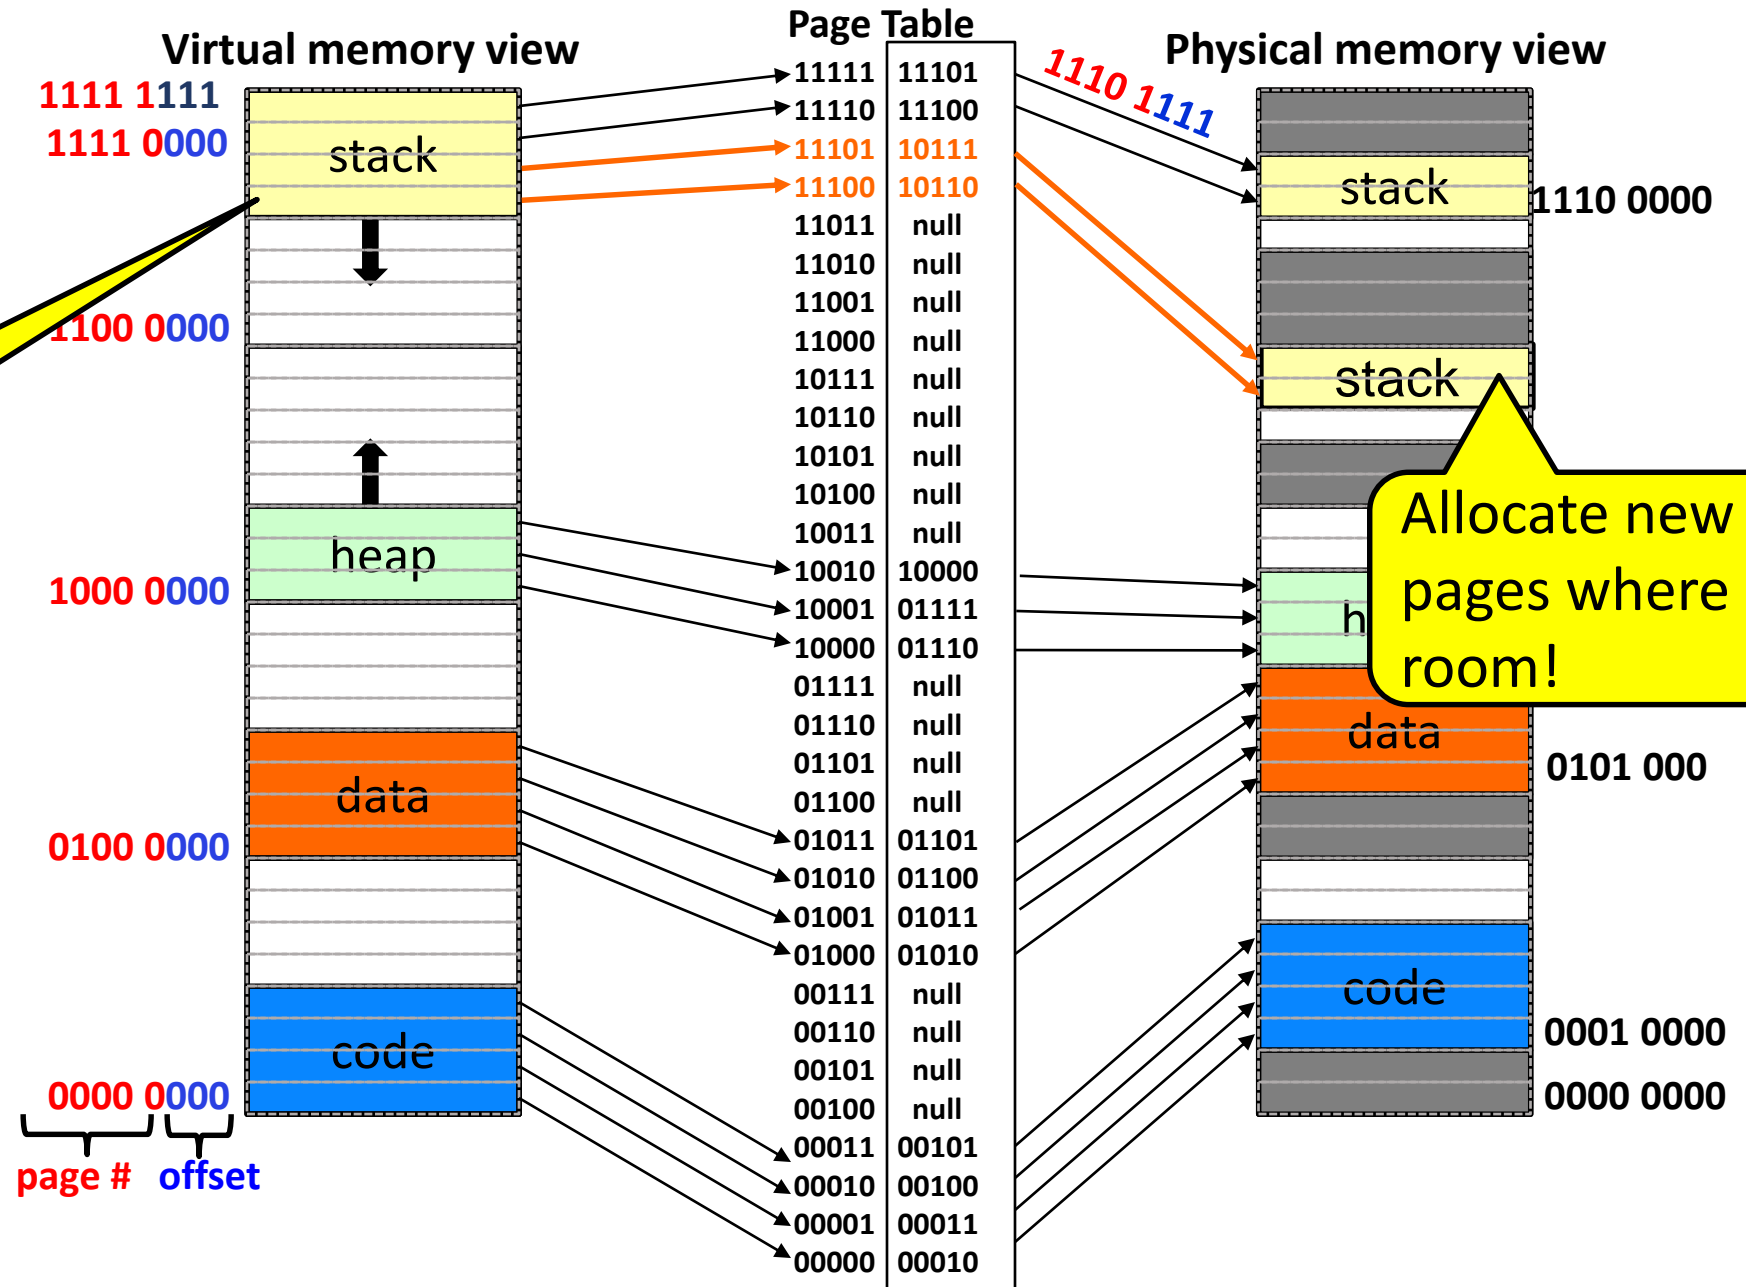
Summary: Paging

What happens
if stack grows
to 1110 0000?



Summary: Paging

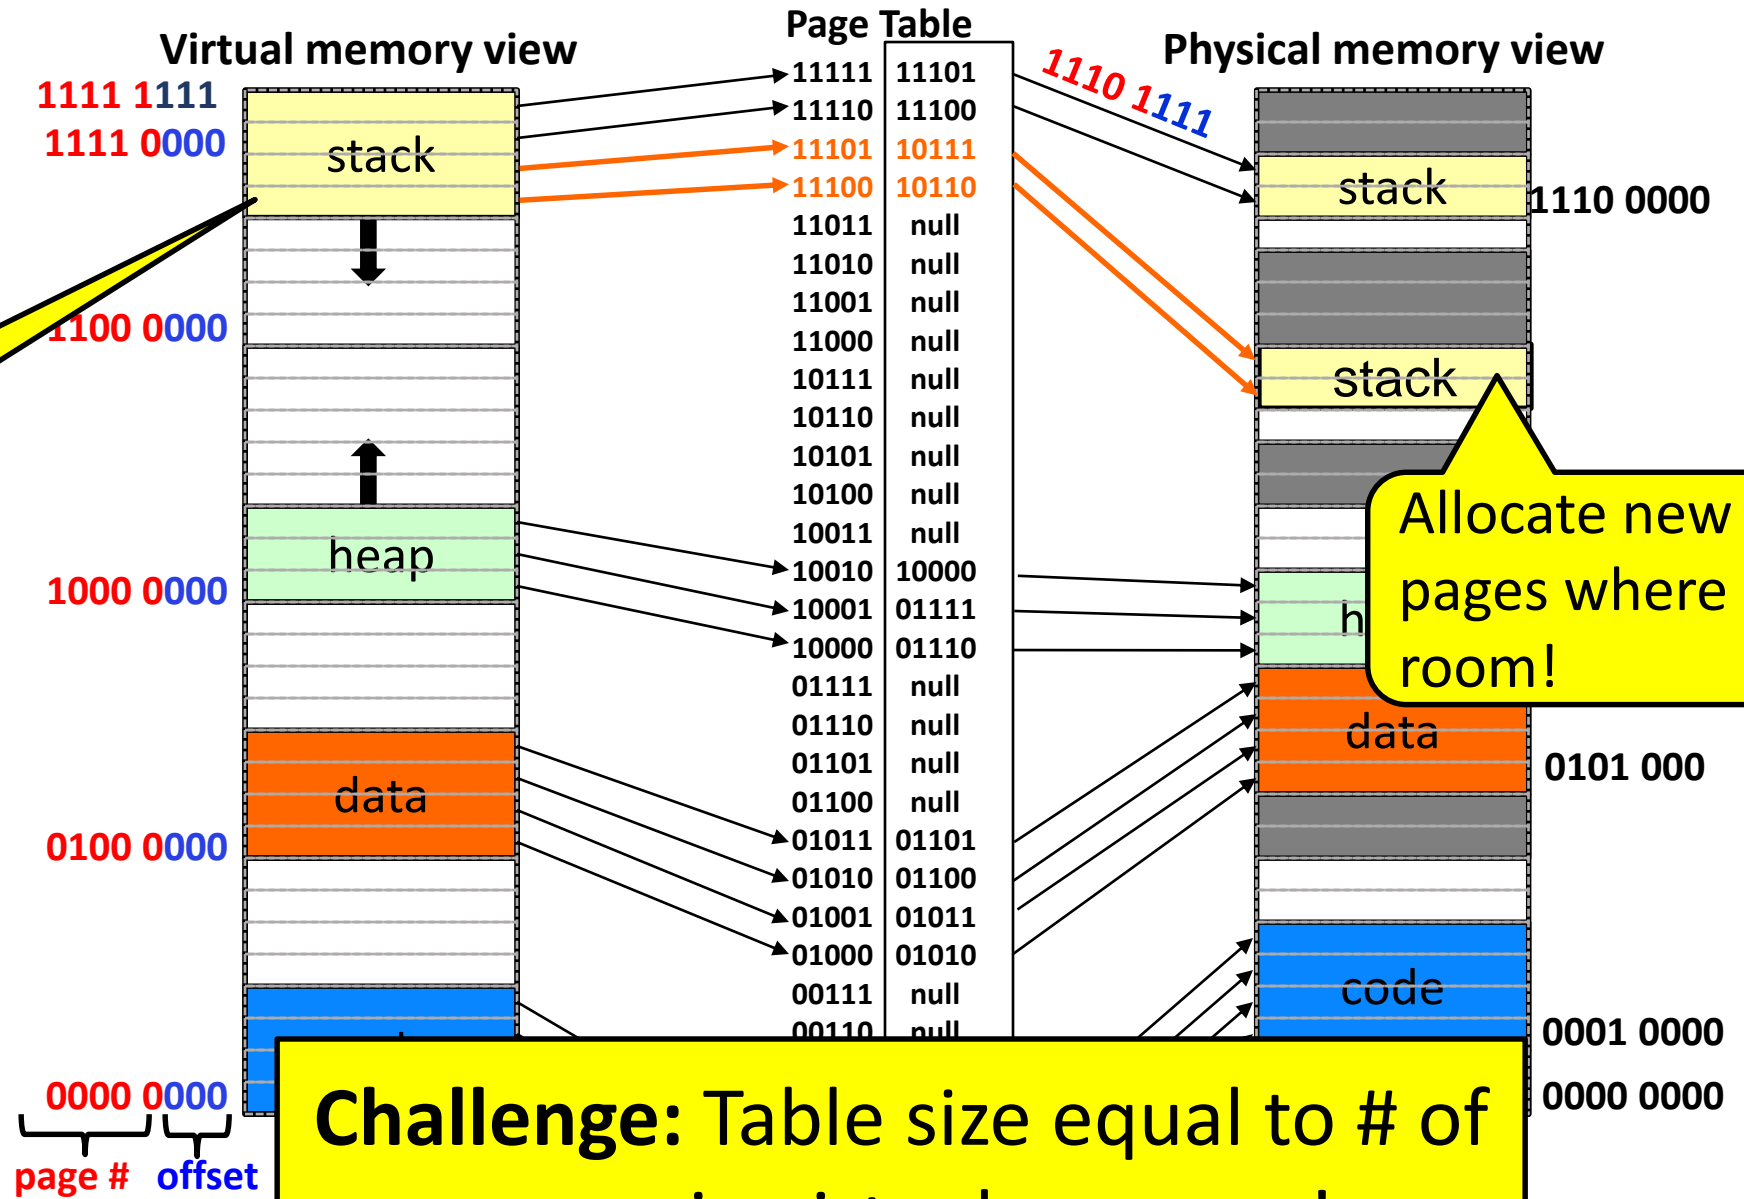
What happens
if stack grows
to 1110 0000?



Allocate new
pages where
room!

Summary: Paging

What happens
if stack grows
to 1110 0000?



How Big is the Page Table?

- Typical page size: 4 KiB
 - How many bits of the address is that? (remember $2^{10} = 1024$)
 - Ans: $4\text{KiB} = 4 \times 2^{10} = 2^{12} \Rightarrow 12$ bits of the address
- So how big is the simple page table for *each* process?
 - $2^{32}/2^{12} = 2^{20}$ (that's about a million entries) x 4 bytes each \Rightarrow 4 MiB
 - When 32-bit machines got started (vax 11/780, intel 80386), this was a lot of memory
- How big is a simple page table on a 64-bit processor (x86_64)?
 - $2^{64}/2^{12} = 2^{52}$ (that's 4.5×10^{15} or 4.5 exa-entries) x 8 bytes each = 36×10^{15} bytes or 36 exa-bytes!!!! This is a ridiculous amount of memory!
- Mostly, the address space is *sparse*, i.e. has holes in it that are not mapped to physical memory
 - So, most of this space is taken up by page tables mapped to nothing

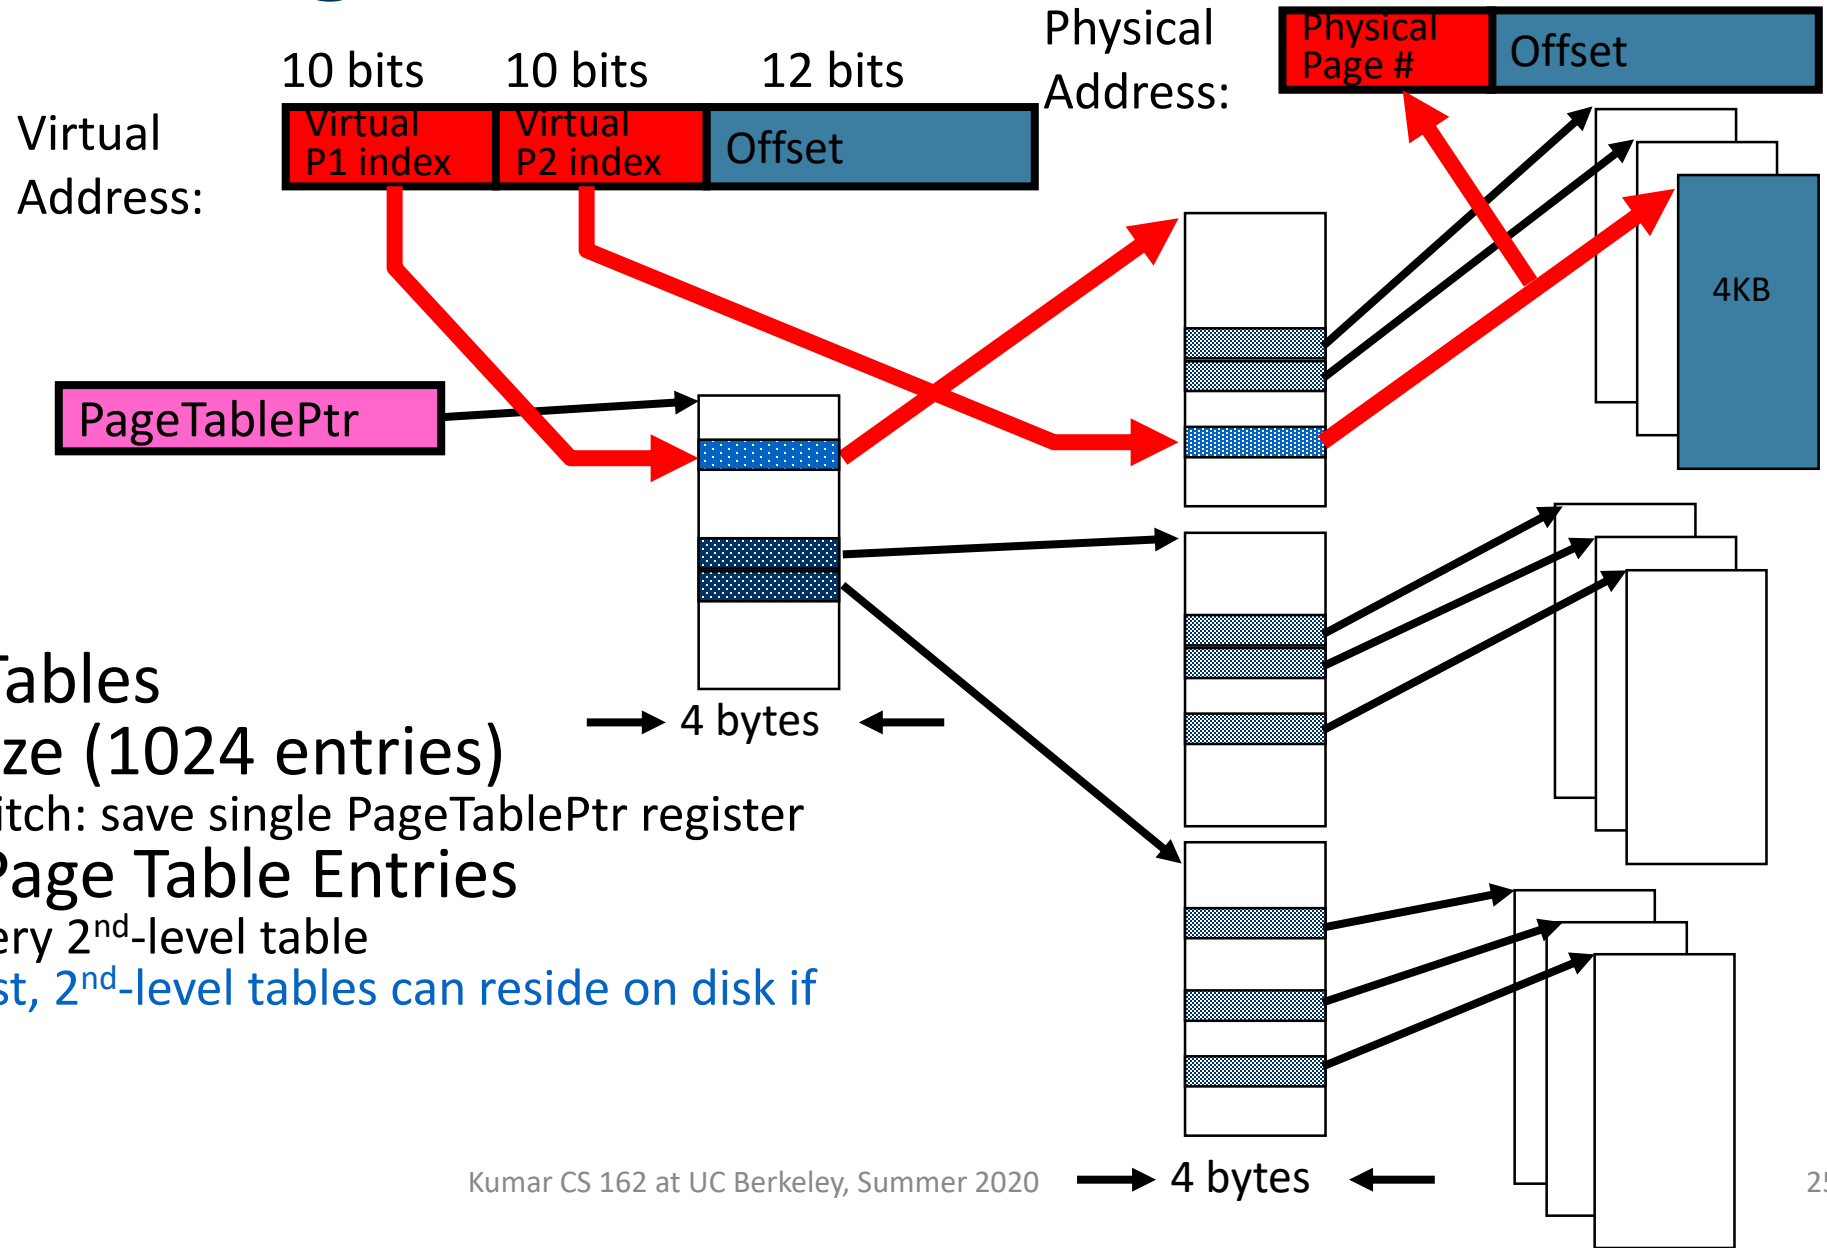
Page Table Discussion

- What provides protection here?
 - Translation (per process) *and* dual-mode operation!
 - Can't let process alter its own page table!
- Analysis
 - Pros
 - Simple memory allocation
 - Easy to share
 - Con: What if address space is sparse?
 - E.g., on UNIX, code starts at 0, stack starts at $(2^{31}-1)$
 - With 1K pages, need 2 million page table entries!
 - Con: What if table really big?
 - Not all pages used all the time \Rightarrow would be nice to have working set of page table in memory
- Simple Page table is way too big!
 - Does it all need to be in memory?
 - How about multi-level paging?
 - or combining paging and segmentation

How to Structure a Page Table

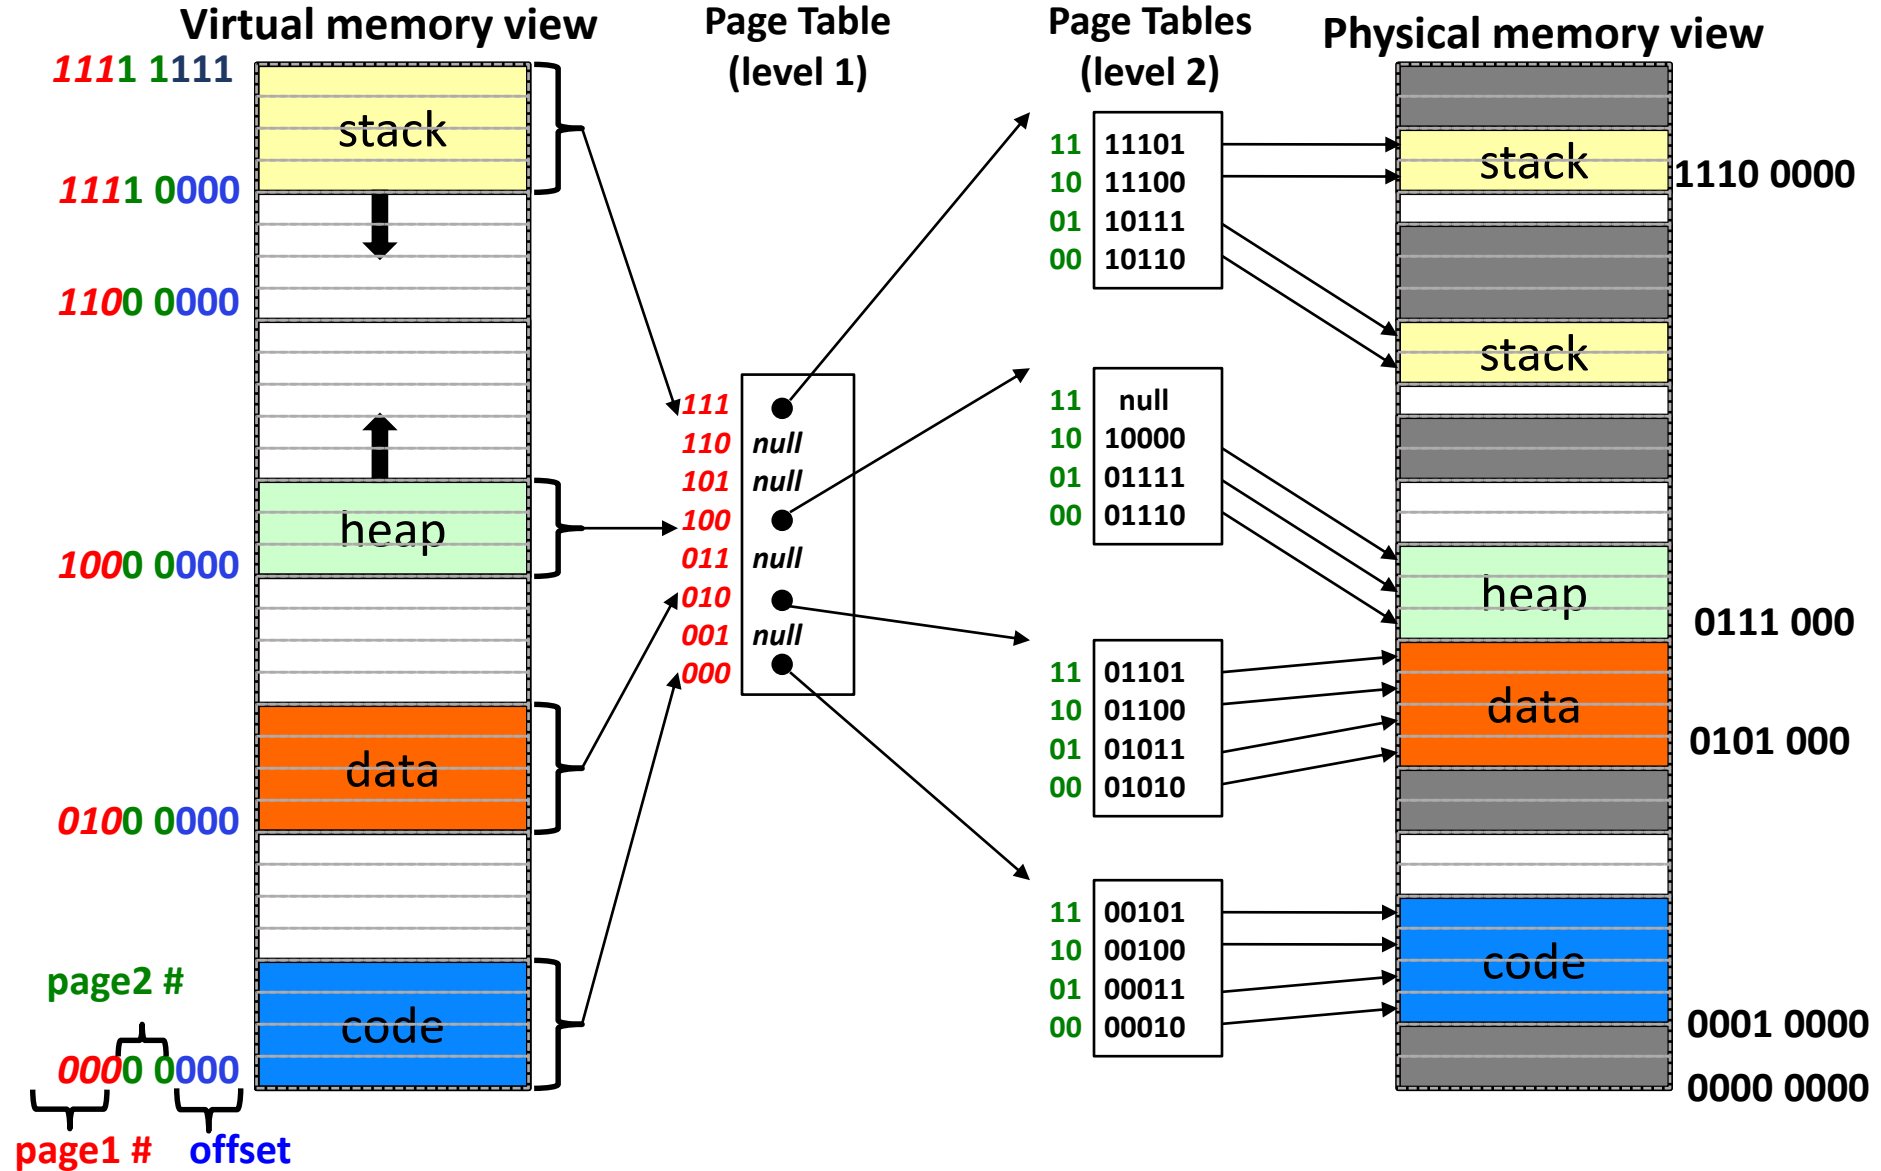
- Page Table is a *map* from VPN to PPN
- Simple page table corresponds to a sparse array
 - VPN is index into table, each entry contains PPN
- What other map structures can you think of?
 - Trees?
 - Hash Tables?

Two-Level Page Table

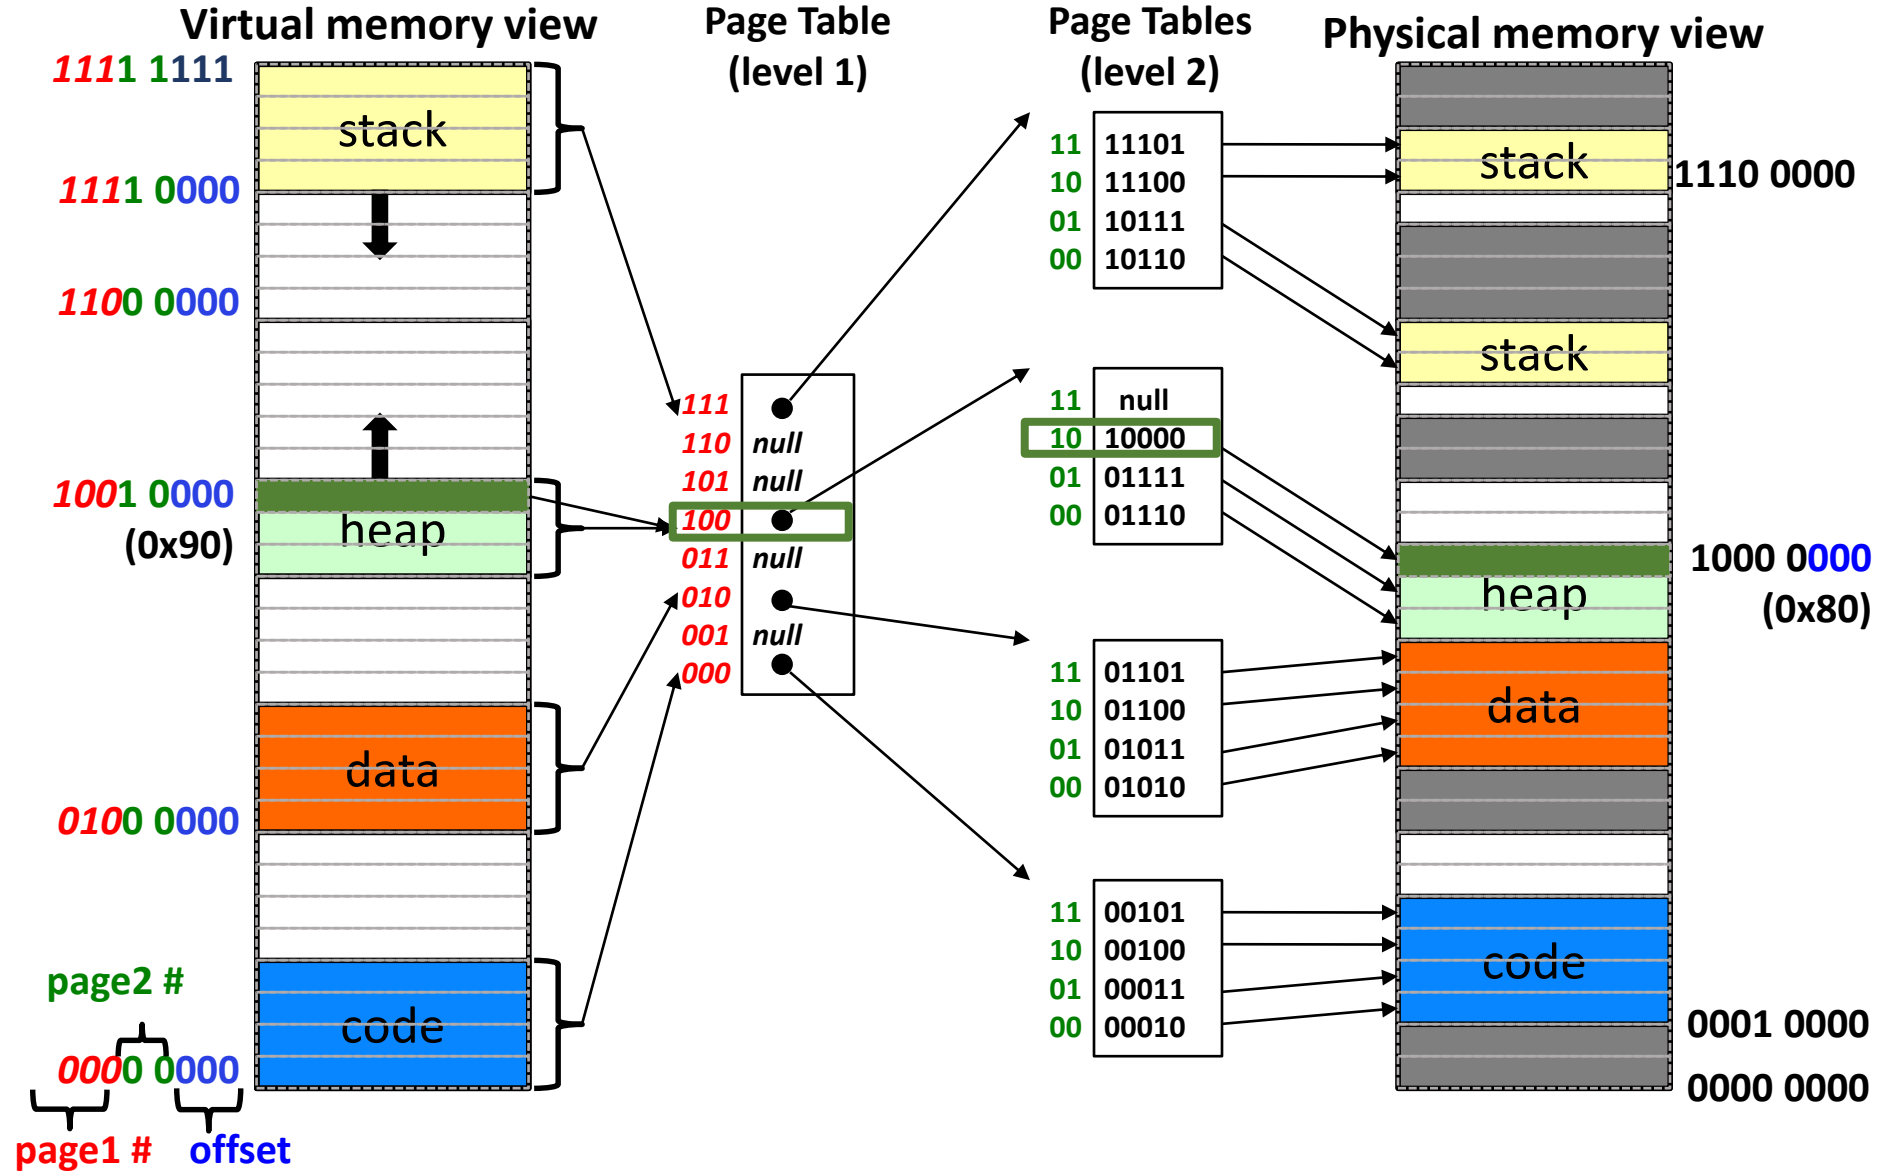


- Tree of Page Tables
- Tables fixed size (1024 entries)
 - On context-switch: save single PageTablePtr register
- Valid bits on Page Table Entries
 - Don't need every 2nd-level table
 - Even when exist, 2nd-level tables can reside on disk if not in use

Summary: Two-Level Paging



Summary: Two-Level Paging



x86 Classic 32-bit Address Translation

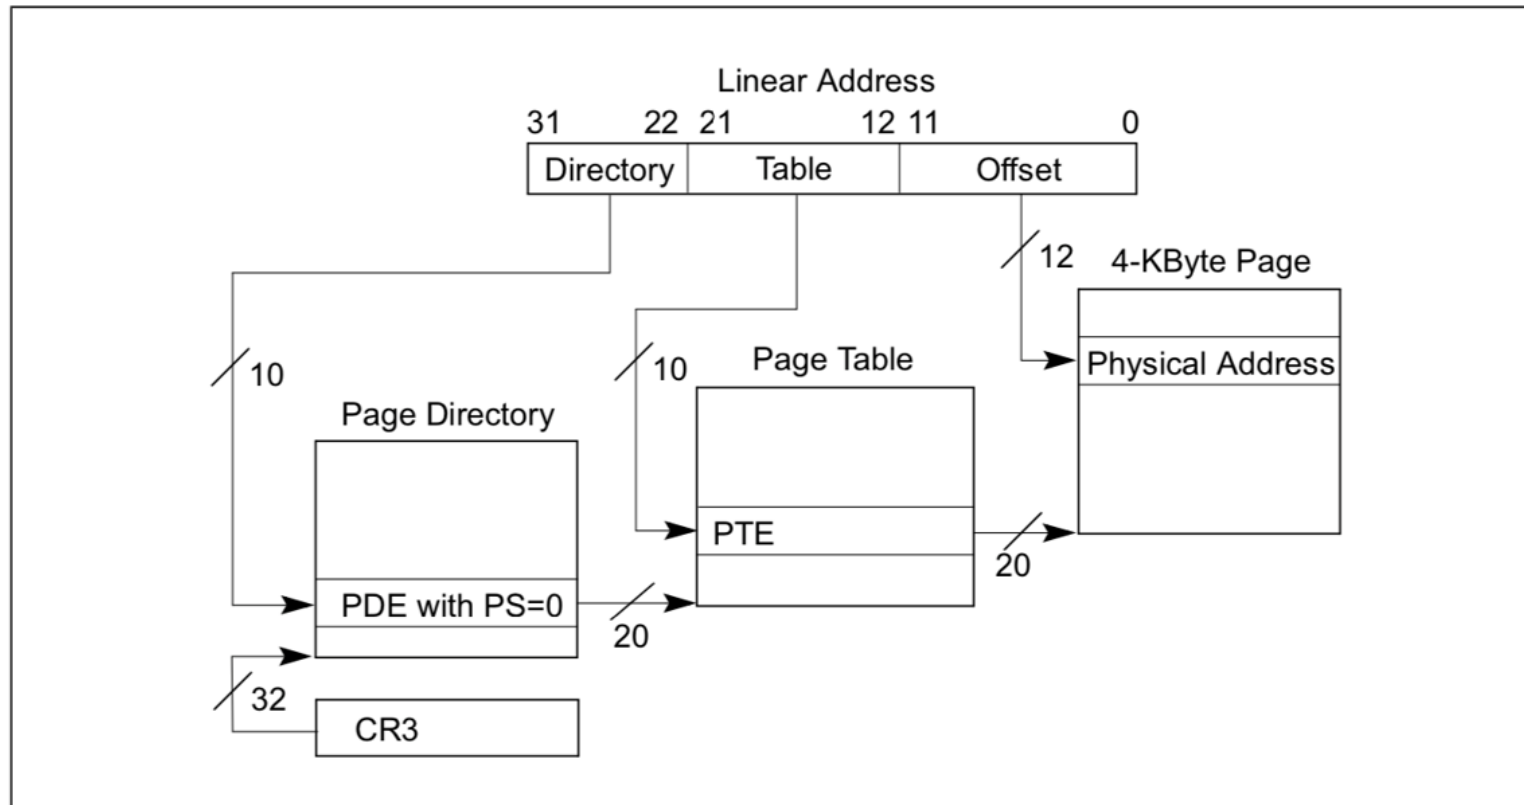
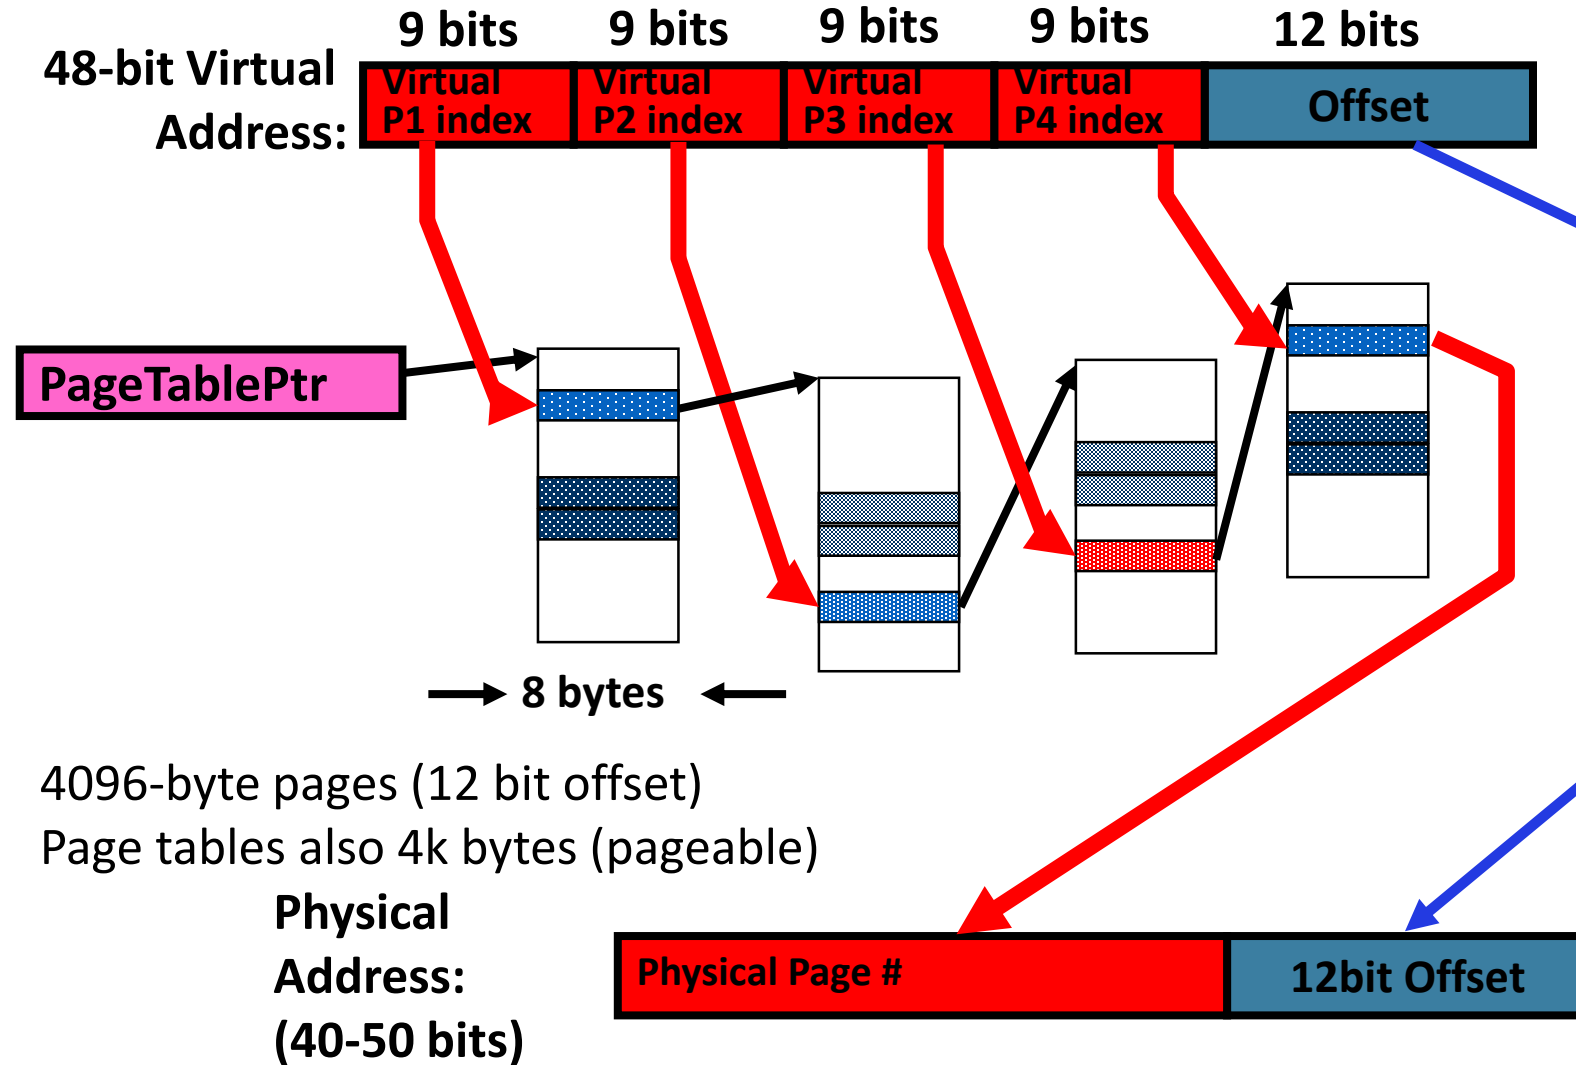


Figure 4-2. Linear-Address Translation to a 4-KByte Page using 32-Bit Paging

x86-64: Four-Level Page Table!



Large 64-bit Address Space

- All current x86-64 processors support 64-bit operations
- 64-bit words but 48-bit addresses

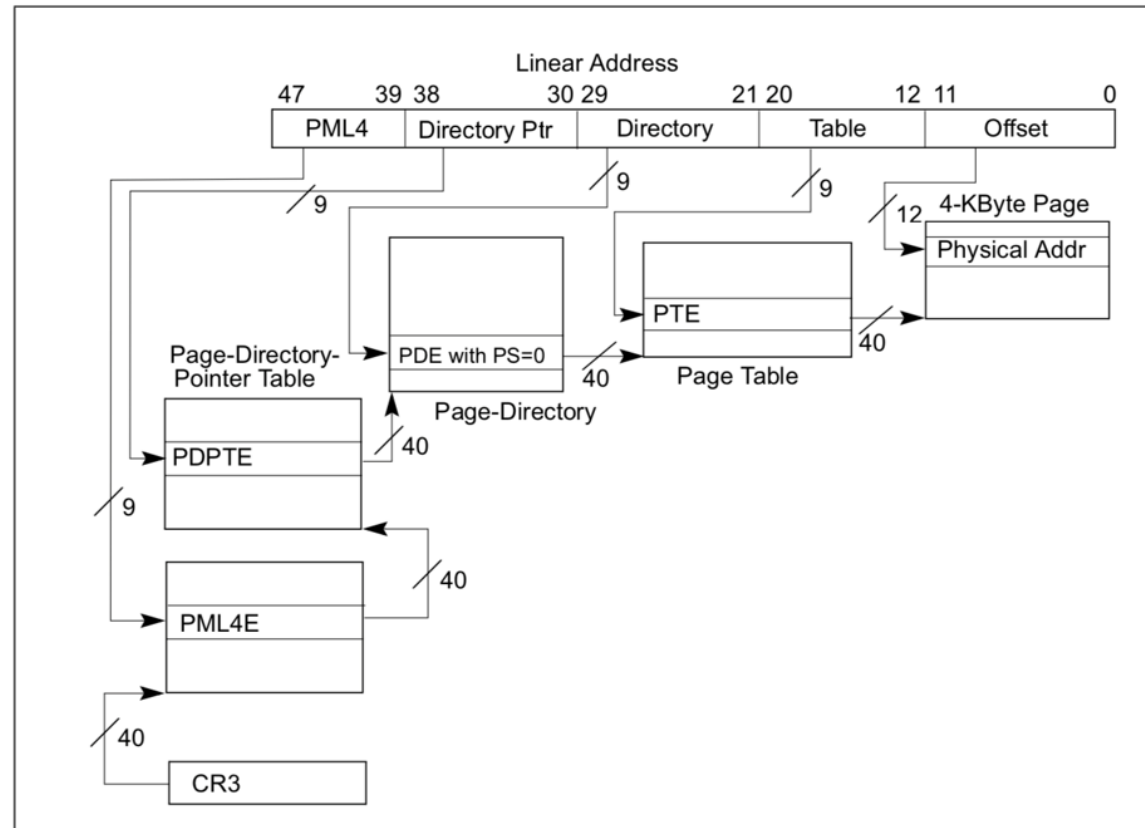


Figure 4-8. Linear-Address Translation to a 4-KByte Page using 4-Level Paging

“Huge Pages” Supported as Well

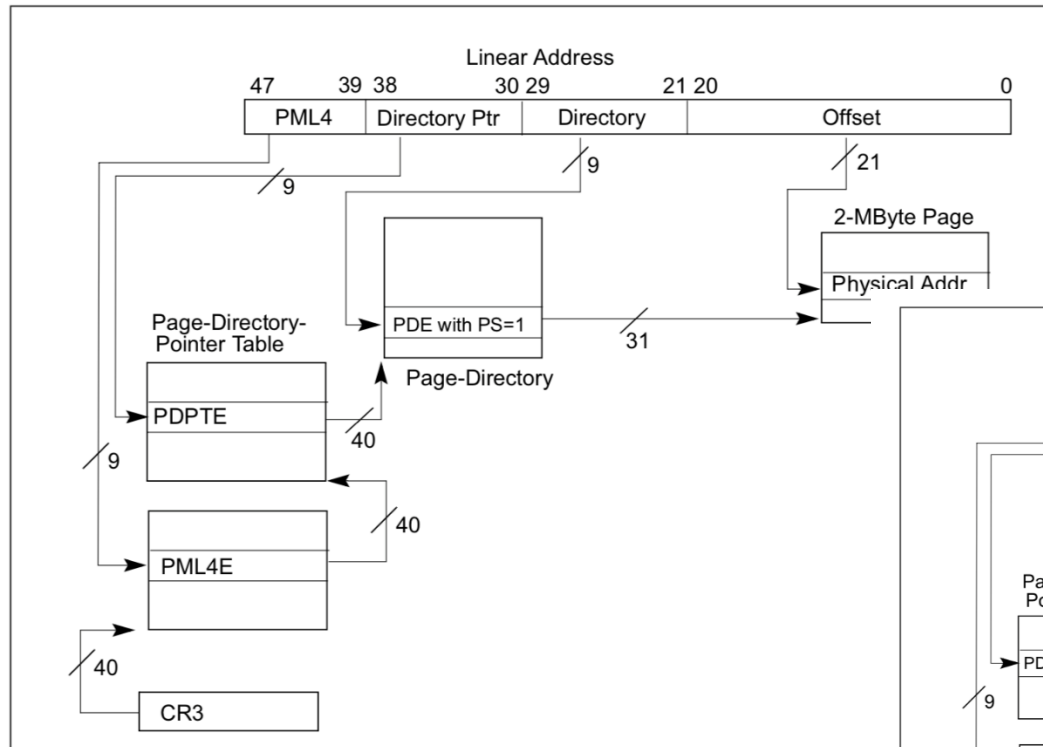


Figure 4-9. Linear-Address Translation to a 2-MByte Page using 4-

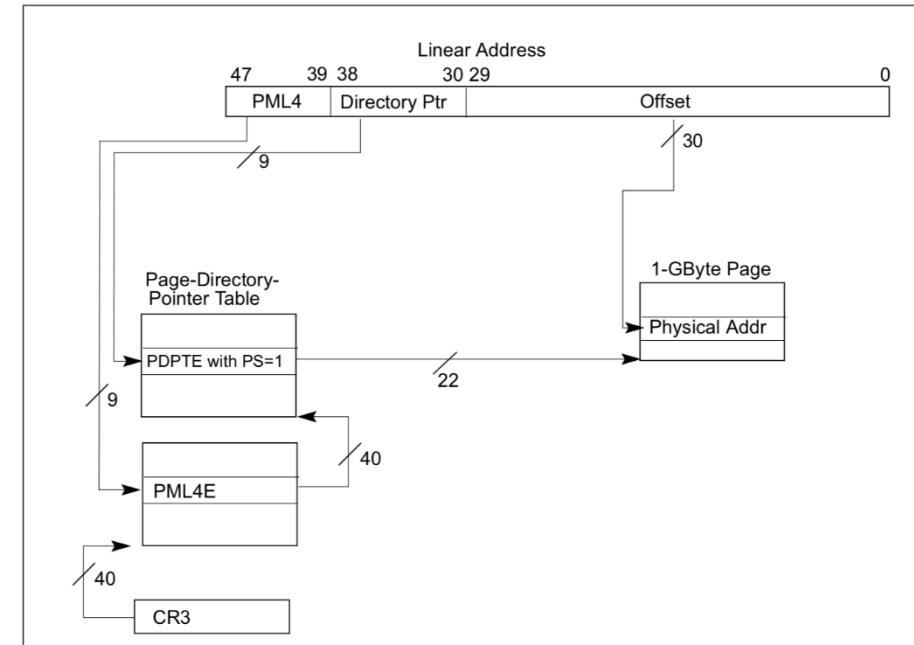


Figure 4-10. Linear-Address Translation to a 1-GByte Page using 4-Level Paging

- Memory is now cheap...

Intel Ice Lake (2019): One More Layer

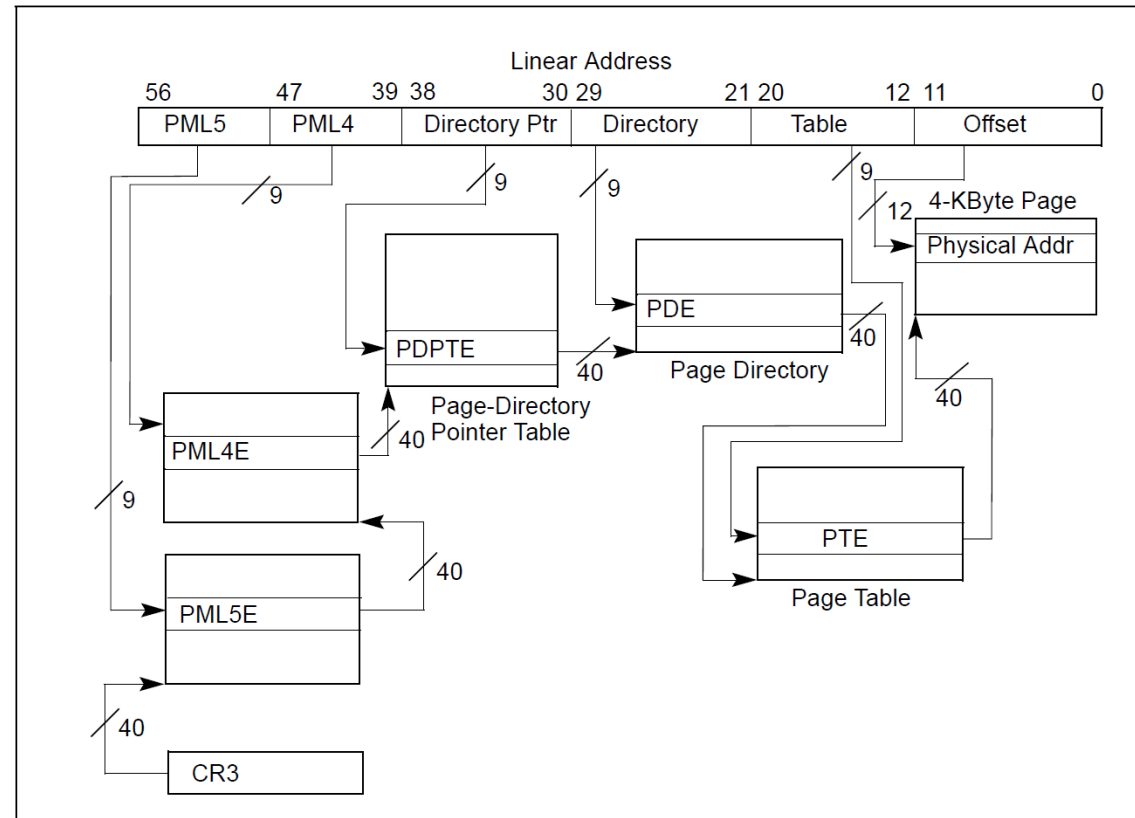
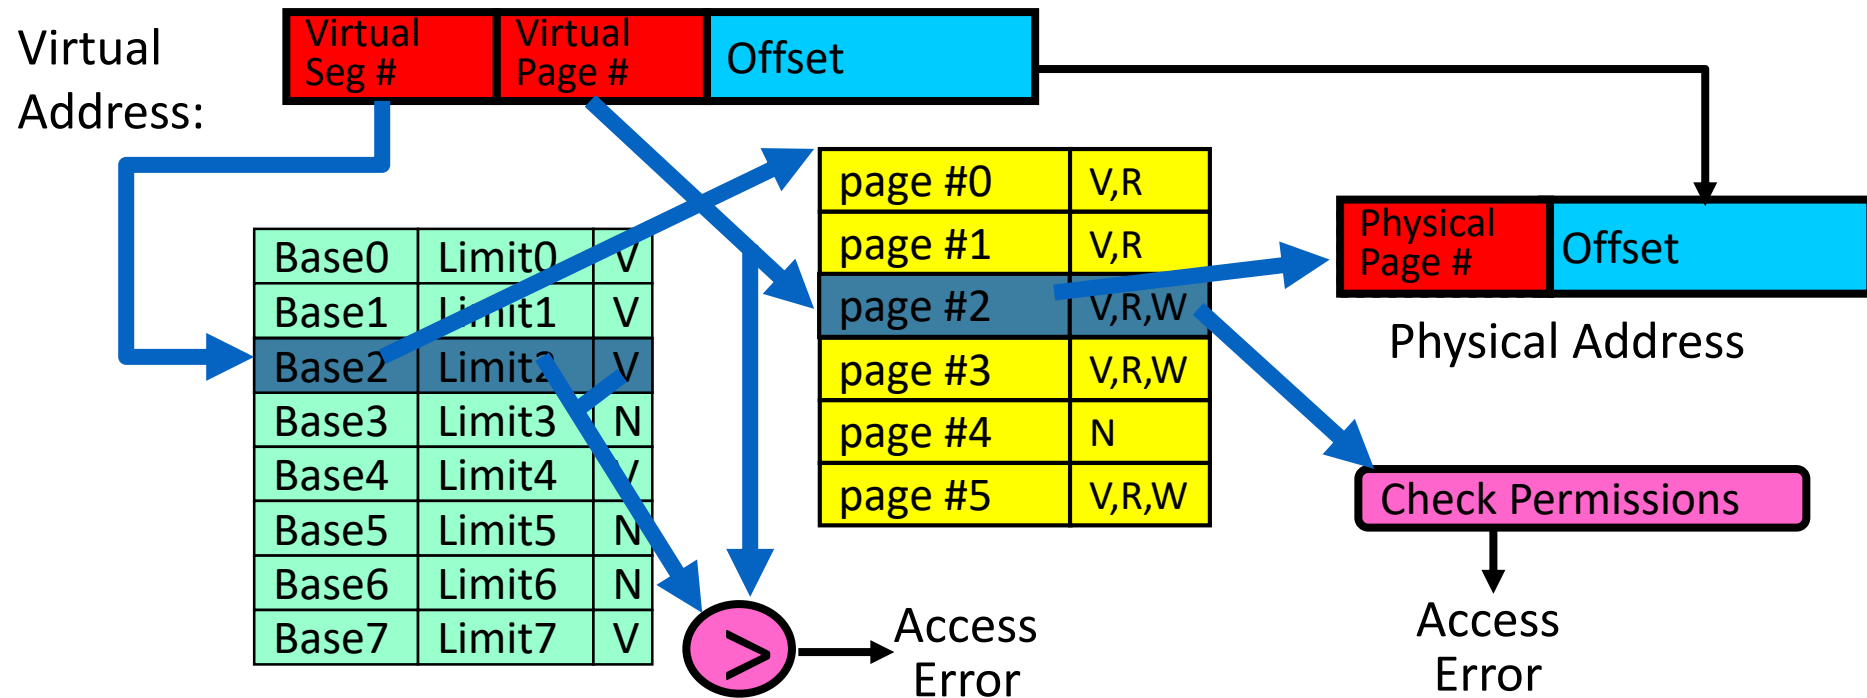


Figure 2-1. Linear-Address Translation Using 5-Level Paging

Multi-Level Translation Analysis

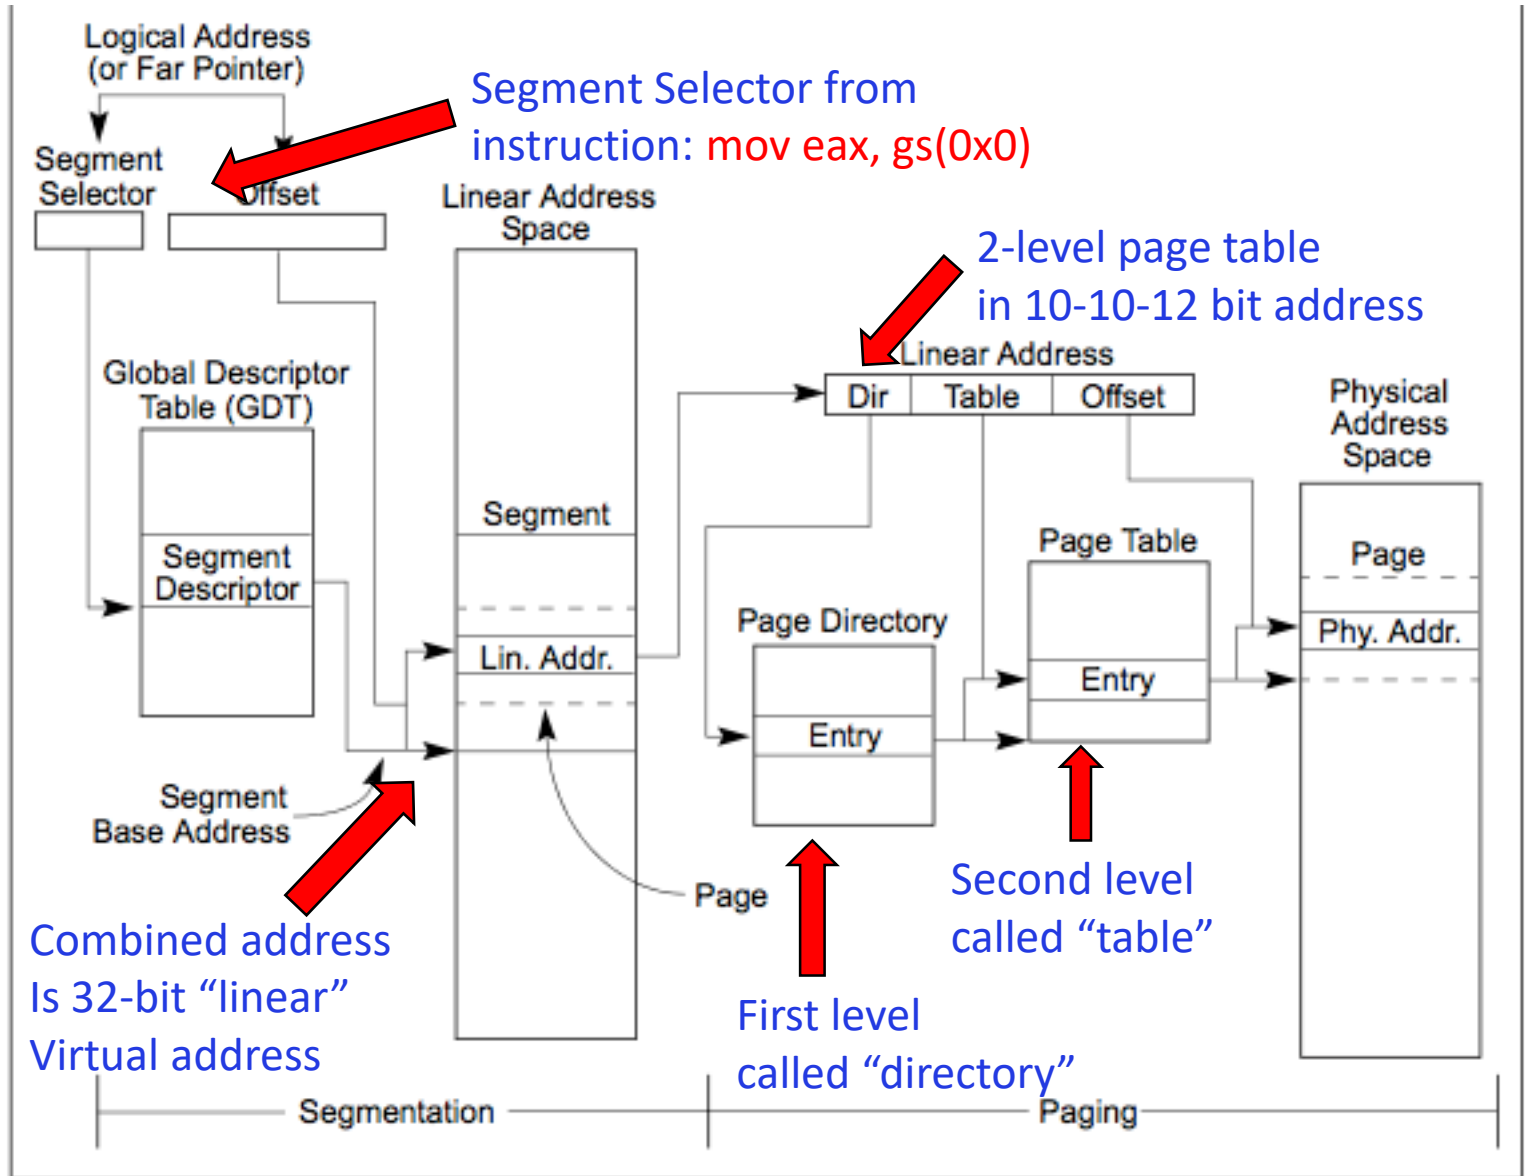
- Pros:
 - Only need to allocate as many page table entries as we need for application
 - In other words, sparse address spaces are easy
 - Easy memory allocation
 - Easy Sharing
 - Share at segment or page level (need additional reference counting)
- Cons:
 - One pointer per page (typically 4K – 16K pages today)
 - Page tables need to be contiguous
 - However, previous example keeps tables to exactly one page in size
 - Two (or more, if >2 levels) lookups per reference
 - Seems very expensive!

Aside: Segments + Pages




- What must be saved/restored on context switch?
 - Contents of top-level segment registers (for this example)
 - Pointer to top-level table (page table)

Aside: x86 Memory with Segmentation



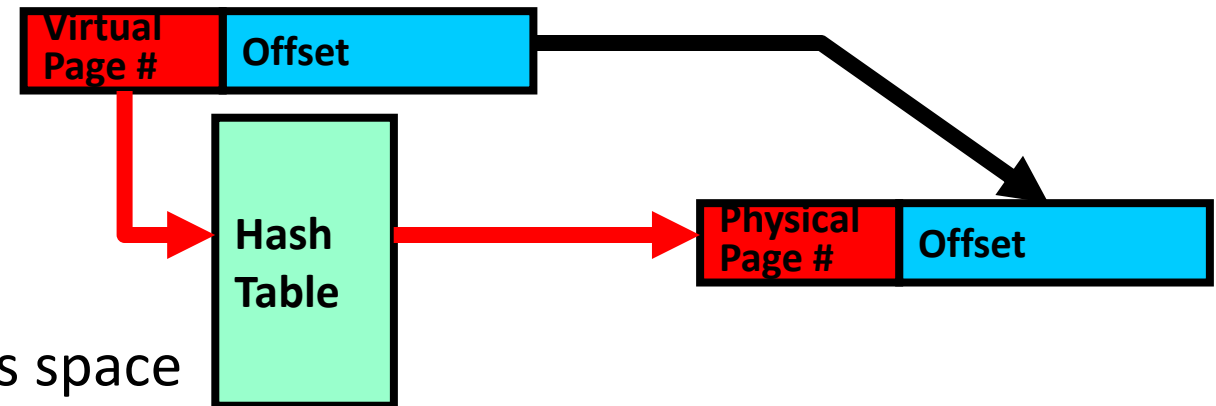
IA-64: 64-bit Address: Six Levels???



Too slow
Too many almost-empty tables

Alternative: Inverted Page Table

- With all previous examples (“Forward Page Tables”)
 - Size of page table is at least proportional to amount of virtual memory allocated to processes
 - Physical memory may be much less
 - Much of process space may be out on disk or not in use
- Answer: use a hash table
 - Called an “Inverted Page Table”
 - Size is independent of virtual address space
 - Directly related to amount of physical memory
 - Very attractive option for 64-bit address spaces
 - PowerPC, UltraSPARC, IA64
- Cons:
 - Complexity of managing hash chains: Often in hardware!
 - Poor cache locality of page table




Address Translation Comparison

	Advantages	Disadvantages
Simple Segmentation	Fast context switching (segment map maintained by CPU)	External fragmentation
Paging (Single-Level)	No external fragmentation Fast and easy allocation	Large table size (~ virtual memory) Internal fragmentation
Paged Segmentation	Table size ~ # of pages in virtual memory Fast and easy allocation	Multiple memory references per page access
Multi-Level Paging		
Inverted Page Table	Table size ~ # of pages in physical memory	Hash function more complex No cache locality of page table

Announcements


- Congrats on finishing Quiz 2!
- Project 2 design doc due tonight
- Homework 4 (Page Walk) comes out tonight (or maybe tomorrow)

How to Translate Addresses Fast Enough?



- The MMU must translate virtual address to physical address on:
 - Every instruction fetch
 - Every load
 - Every store
- More than one translation for EVERY instruction
 - Each one requires a page table ***tree traversal*** (!)
 - How to simplify this???

Where and What is the MMU?

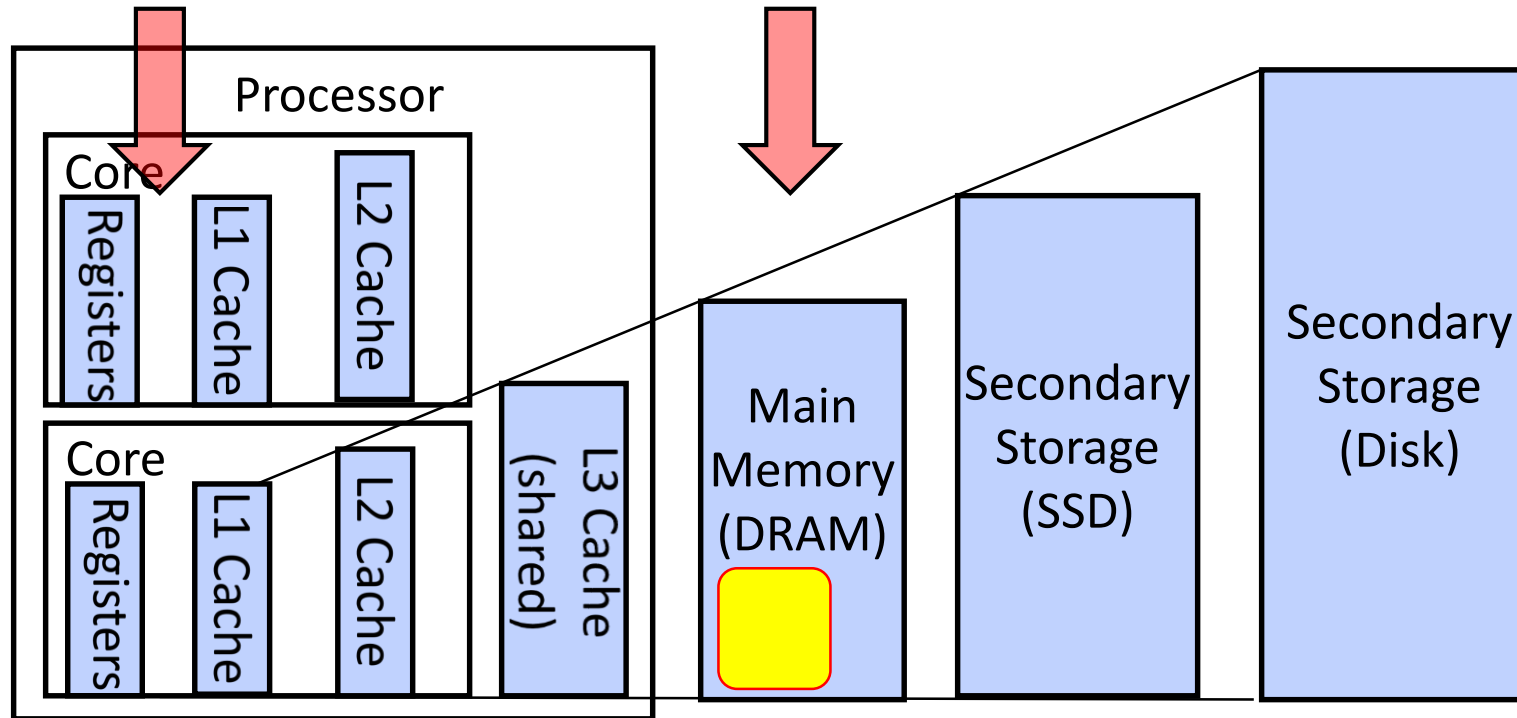


- On every memory reference (I-fetch, Load, Store), MMU reads (multiple levels of) page table entries to get physical frame or FAULT
 - Through the caches to the memory
 - Then read/write the physical location

Recall: Memory Hierarchy

Address Translation
needs to occur here

Page table lives here
(perhaps cached)



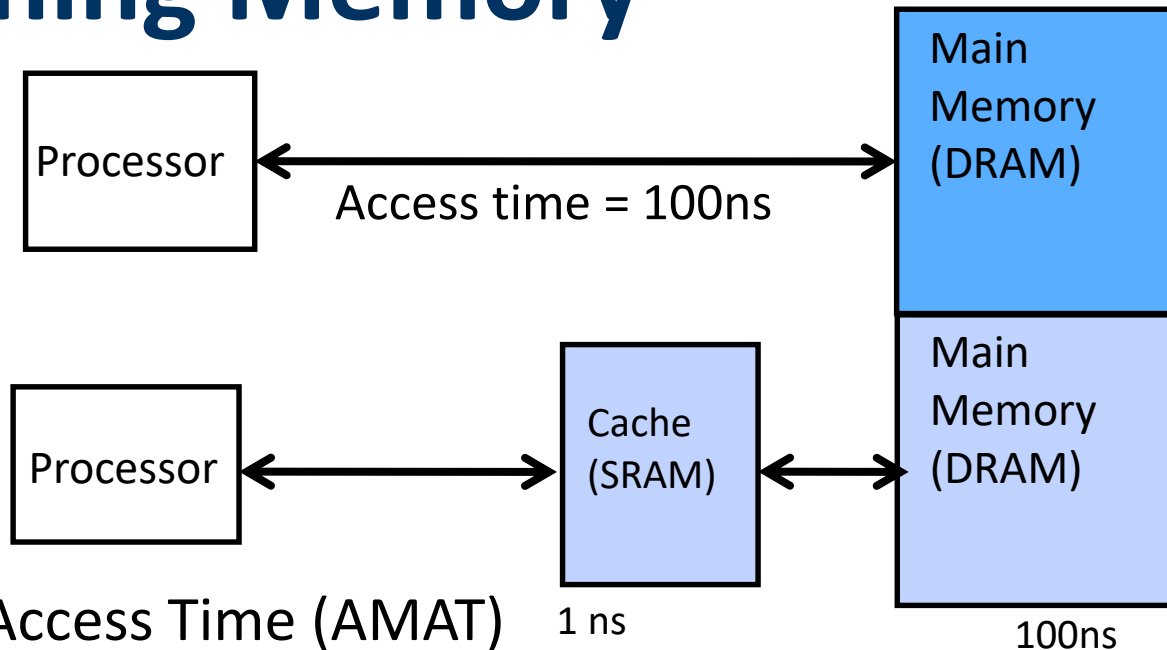
Speed (ns):	0.3	1	3	10-30	100	100,000 (0.1 ms)	10,000,000 (10 ms)
Size (bytes):	100Bs	10kB	100kB	MBs	GBs	100GBs	TBs

- Large memories are slow
- Small memories are fast

Recall: Caches

- **Cache**: a repository for copies that can be accessed more quickly than the original
 - Make frequent case fast and infrequent case less dominant
- Caching underlies many techniques used to make computers fast
 - Can cache: memory locations, address translations, pages, file blocks, file names, network routes, etc...
- Key measure: Average Access time =
$$(\text{Hit Rate} \times \text{Hit Time}) + (\text{Miss Rate} \times \text{Miss Time})$$

Recall: Caching Memory




- Average Memory Access Time (AMAT)
- $$= (\text{Hit Rate} \times \text{HitTime}) + (\text{Miss Rate} \times \text{MissTime})$$
- Where $\text{HitRate} + \text{MissRate} = 1$
- $\text{HitRate} = 90\% \Rightarrow \text{AMAT} = (0.9 \times 1) + (0.1 \times 101) = 11.1 \text{ ns}$
- $\text{HitRate} = 99\% \Rightarrow \text{AMAT} = (0.99 \times 1) + (0.01 \times 101) = 2.01 \text{ ns}$
- MissTime_{L1} includes $\text{HitTime}_{L1} + \text{MissPenalty}_{L1} \equiv \text{HitTime}_{L1} + \text{AMAT}_{L2}$

Recall: Why Does Caching Help? Locality!


- **Temporal Locality** (Locality in Time)

- Keep recently accessed data items closer to processor




- **Spatial Locality** (Locality in Space)

- Move contiguous blocks to the upper levels

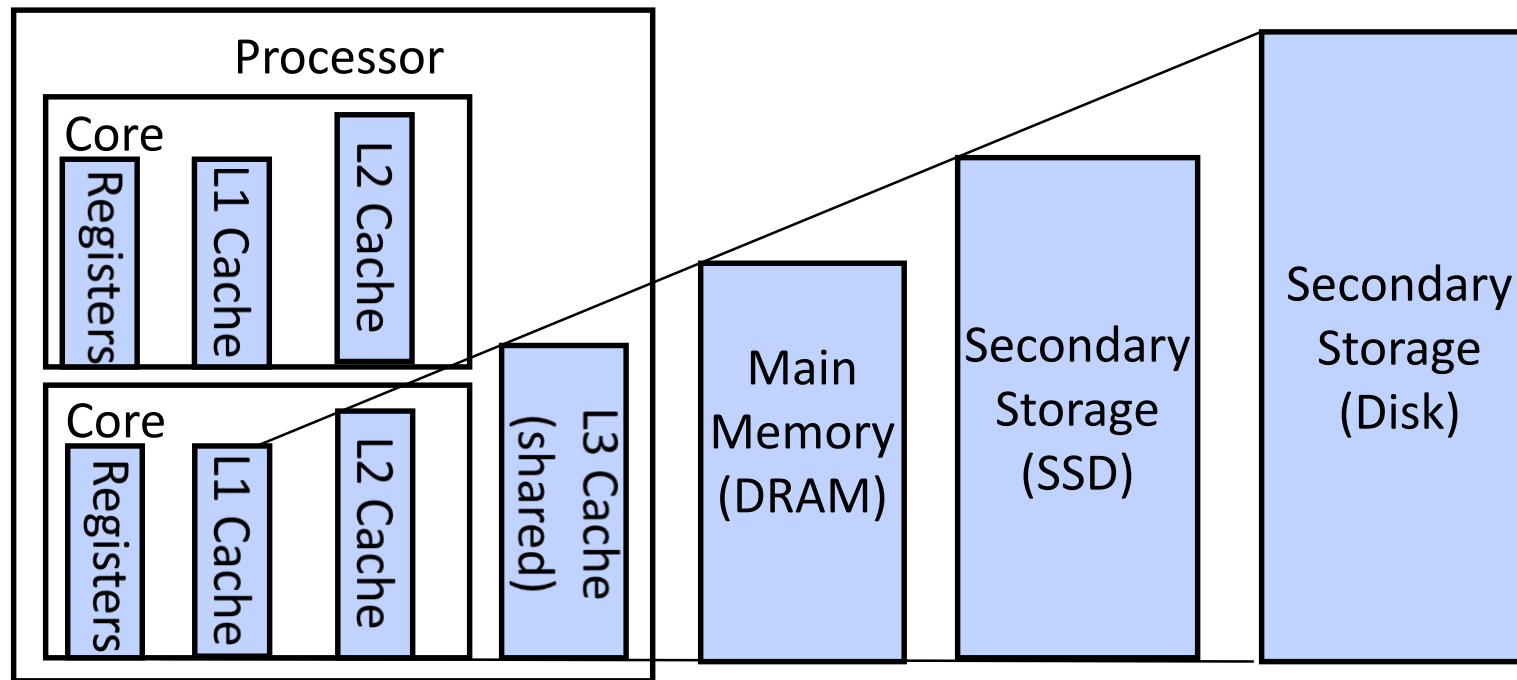


Recall: Working Set Model

- As a program executes it transitions through a sequence of “working sets” consisting of varying sized subsets of the address space



Recall: Memory Hierarchy

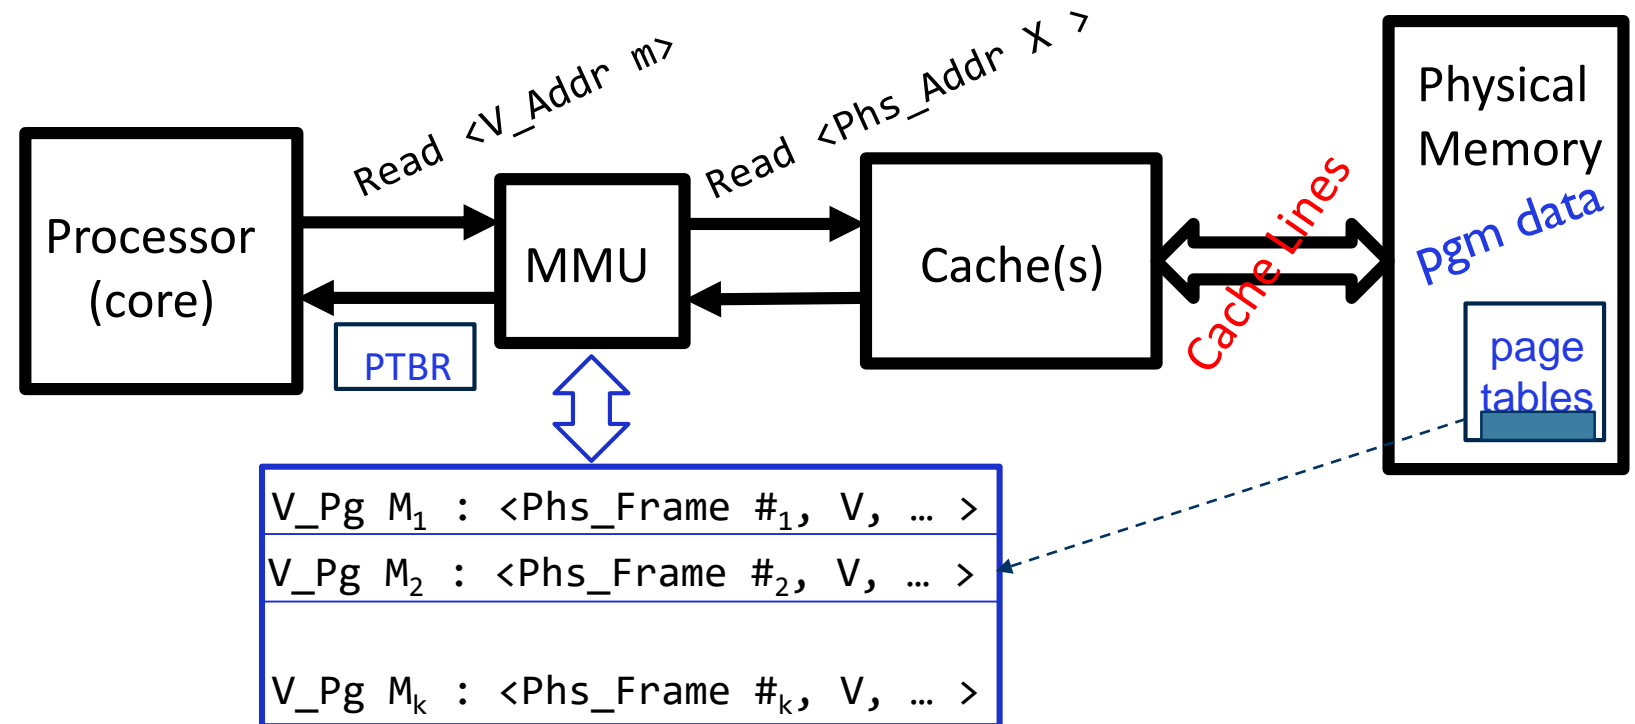


Speed (ns):	0.3	1	3	10-30	100	100,000 (0.1 ms)	10,000,000 (10 ms)
Size (bytes):	100Bs	10kB	100kB	MBs	GBs	100GBs	TBs

- Take advantage of the principle of locality to:
 - Present as much memory as in the cheapest technology
 - Provide access at speed offered by the fastest technology

Making Address Translation Fast

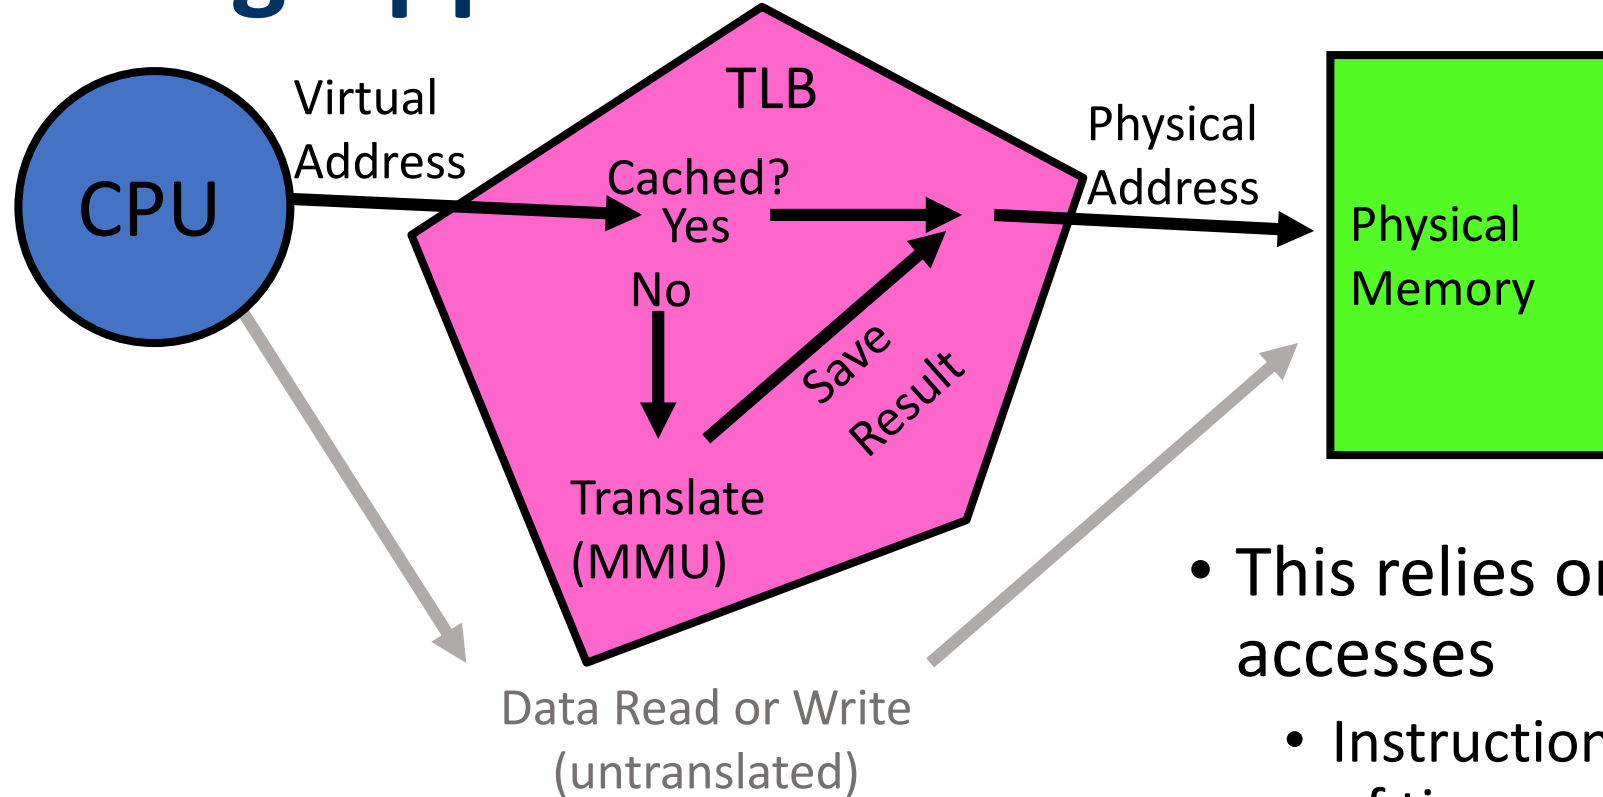
- Cache results of recent translations
 - Separate from memory cache
 - Cache PTEs using Virtual Page Number as the key



Translation Lookaside Buffer (TLB)

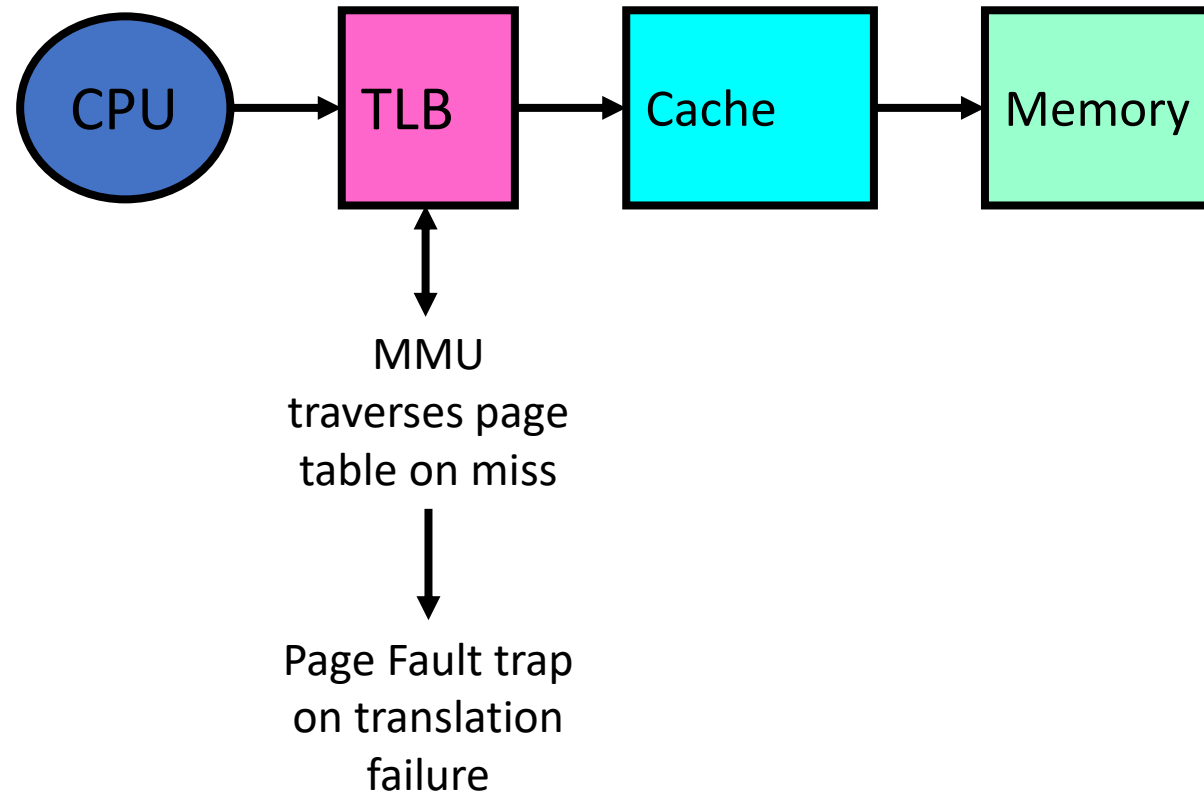
- Record recent Virtual Page # to Physical Frame # translations
- If present in the TLB, can translate address without reading page table
 - Caches the end-to-end result, even if the translation involved multiple levels
- Was invented by Sir Maurice Wilkes – *prior to caches*
- People realized “if it’s good for page tables, why not the rest of the data in memory?”
- On a *TLB miss*, the page tables may be cached, so only go to memory when both miss

Caching Applied to Address Translation




- This relies on locality of PTE accesses
 - Instruction accesses spend a lot of time on the same page
 - Stack accesses have locality
 - Data accesses???

The Big Picture



What Kind of Cache for TLB?



- Remember all those cache design parameters and trade-offs?
- Amount of Data = $N * L * K$
- Write Policy (write-thru, write-back), Eviction Policy (LRU, ...)


How might organization of a TLB differ from that of a conventional instruction or data cache?

Let's do some review...

Recall: Sources of Cache Misses

- **Compulsory** (cold start or first reference): first access to a block
 - “Cold” fact of life: not a whole lot you can do about it
 - Note: If you are going to run “billions” of instruction, Compulsory Misses are insignificant
- **Capacity**:
 - Cache cannot contain all blocks access by the program
 - Solution: increase cache size
- **Conflict** (collision):
 - Multiple memory locations mapped to the same cache location
 - Solution 1: increase cache size
 - Solution 2: increase associativity
- **Coherence** (Invalidation): other process (e.g., I/O) updates memory

Recall: Finding a Block in a Cache?



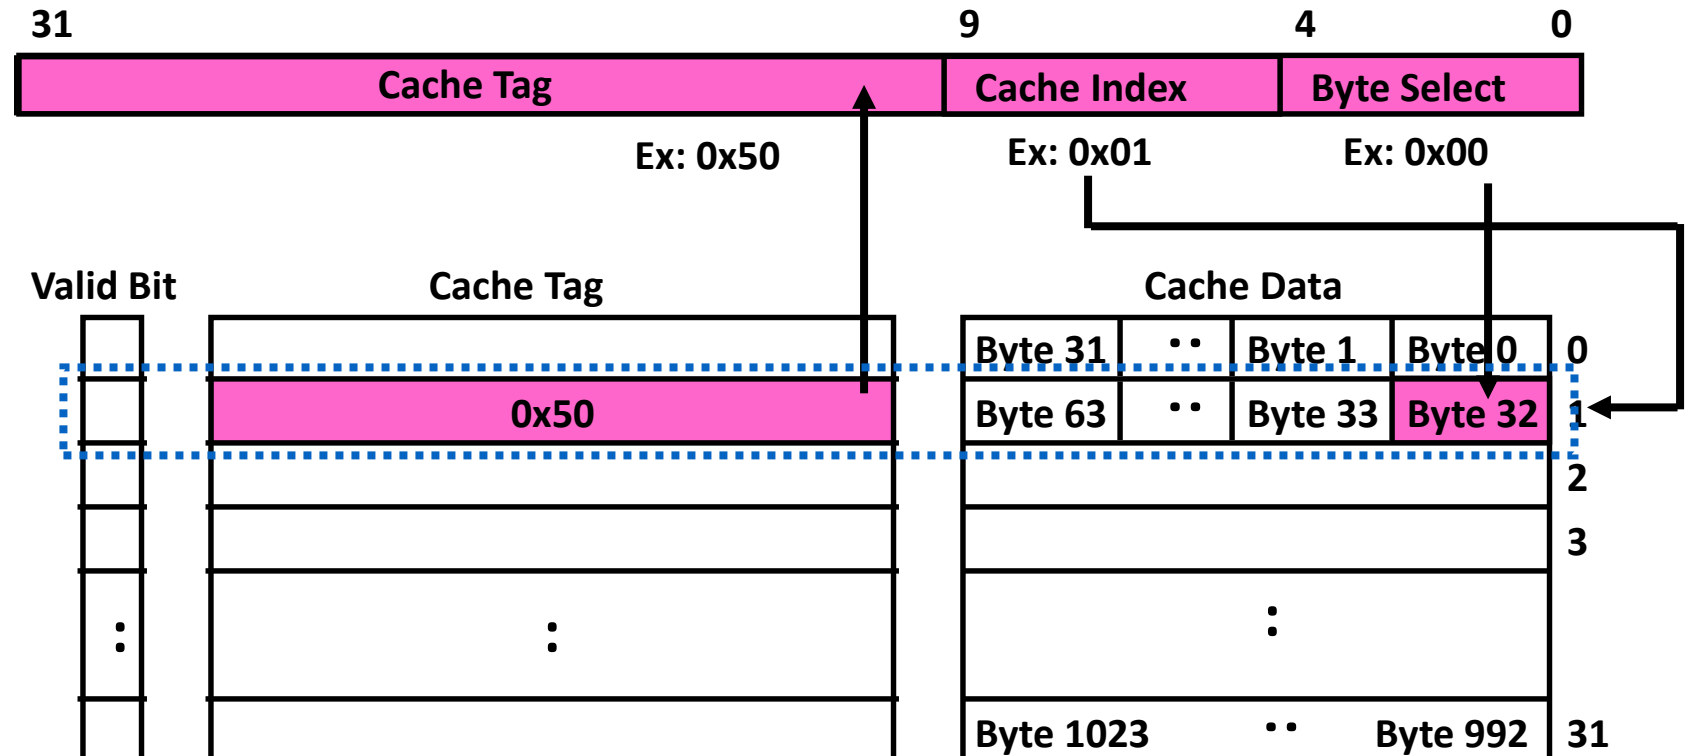
- **Block** is minimum quantum of caching
 - Data select field used to select data within block
 - Many caching applications don't have data select field
- **Tag** used to identify actual the block (what address?)
 - If no candidates match, then declare cache miss
- **Index** Used to Lookup Candidates in Cache
 - Index identifies the set of possibilities (check tag)

Recall: Direct-Mapped Cache

- Direct Mapped 2^N byte cache:
 - The uppermost $(32 - N)$ bits are always the Cache Tag
 - The lowest M bits are the Byte Select (Block Size = 2^M)

Example: 1 KB Direct Mapped Cache with 32 B Blocks

- Index chooses potential block
- Tag checked to verify block
- Byte select chooses byte within block

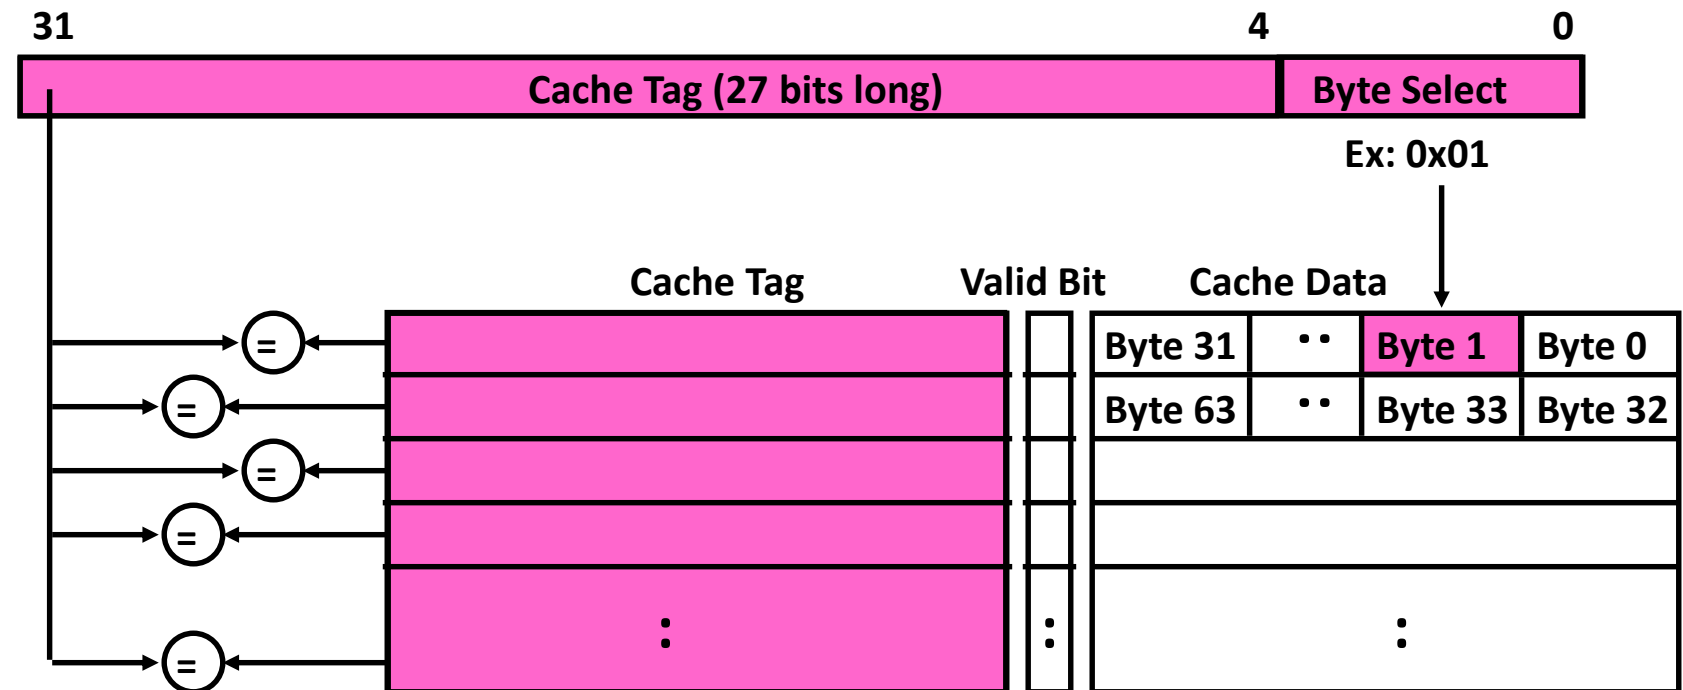


Recall: Fully Associative Cache

- **Fully Associative:** Every block can hold any line
 - Address does not include a cache index
 - Compare tags of all cache entries in parallel

Example: Block Size=32B blocks

- We need N 27-bit comparators
- Still have byte select to choose from within block

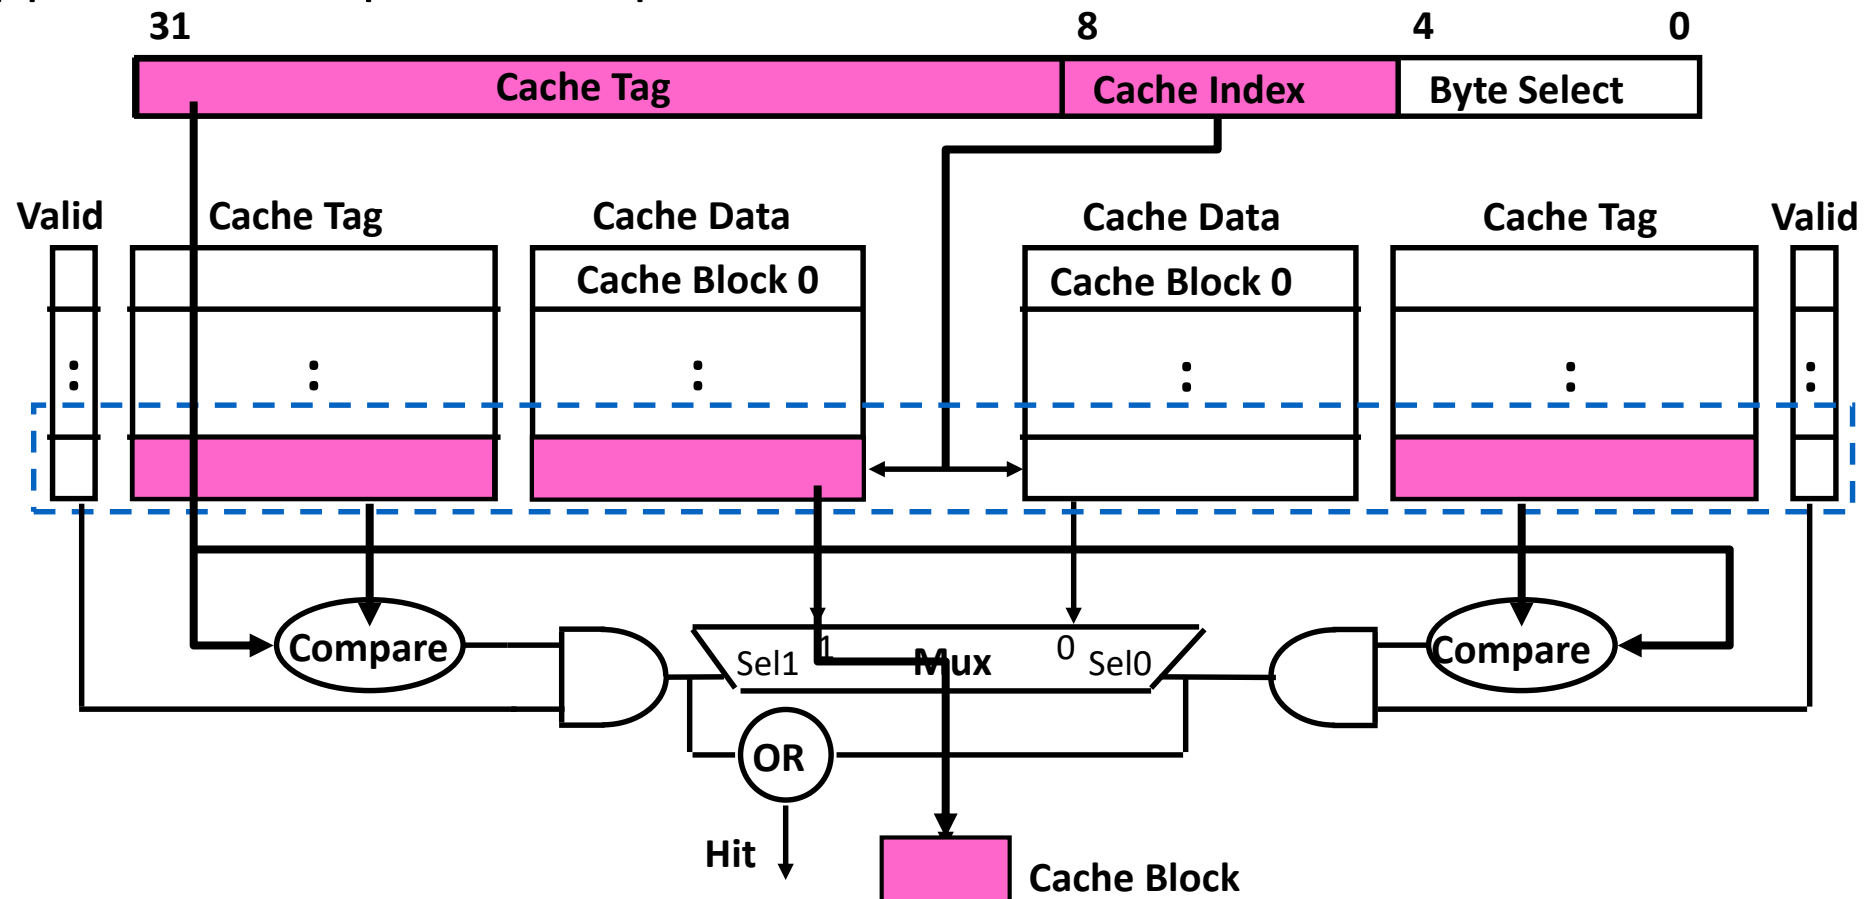


Recall: Set-Associative Cache

- **N-way set associative**: N entries per Cache Index
 - N direct mapped caches operates in parallel

Example: Two-way set associative cache


- Cache Index selects a “set” from the cache
- Two tags in the set are compared to input in parallel
- Data is selected based on the tag result



Where Does a Block Get Placed in a Cache?

- Example: Block 12 placed in 8-block cache
 - Address space has 32 blocks


32-Block Address Space:



Block no.	0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	0	1
											1	1	1	1	1	1	1	1	1	2	2	2	2	2	2	2	2	2	3	3	3	


Direct mapped:

block 12 can go
only into block 4
($12 \bmod 8$)




Set associative:

block 12 can go
anywhere in set 0 ($12 \bmod 4$)



Fully associative:

block 12 can go
anywhere



Recall: Which Block to Replace on a Miss?

- Easy for Direct Mapped: Only one possibility
- Set Associative or Fully Associative:
 - Random
 - LRU (Least Recently Used)
- Miss rates for a workload:

Size	2-way		4-way		8-way	
	LRU	Random	LRU	Random	LRU	Random
16 KB	5.2%	5.7%	4.7%	5.3%	4.4%	5.0%
64 KB	1.9%	2.0%	1.5%	1.7%	1.4%	1.5%
256 KB	1.15%	1.17%	1.13%	1.13%	1.12%	1.12%

Recall: What Happens on a Write?

- **Write through**: The information is written to both the block in the cache and to the block in the lower-level memory
- **Write back**: The information is written only to the block in the cache
 - Modified cache block is written to main memory only when it is replaced
 - Question is block clean or dirty?
- Pros and Cons of each?
 - WT:
 - PRO: read misses cannot result in writes
 - CON: Processor held up on writes unless writes buffered
 - WB:
 - PRO: repeated writes not sent to DRAM
processor not held up on writes
 - CON: More complex
Read miss may require writeback of dirty data

What does our understanding of caches tell us about TLB design?

What TLB Organization Makes Sense?

- Needs to be really fast
 - Seems to argue for Direct Mapped or Low Associativity
- However, needs to have very few conflicts!
 - With TLB, the Miss Time extremely high! (PT traversal)
 - Cost of Conflict (Miss Time) is high
 - Hit Time – dictated by clock cycle
- **Thrashing:** continuous conflicts between accesses
 - What if use low order bits of page as index into TLB?
 - First page of code, data, stack may map to same entry
 - Need 3-way associativity at least?
 - What if use high order bits as index?
 - TLB mostly unused for small programs


TLB Organization: Include Protection

- How big does TLB actually have to be?
 - Usually fewer entries than the cache (why?)
 - Not very big, can support higher associativity
- Small TLBs usually organized as fully-associative cache
 - Lookup is by Virtual Address
 - Returns Physical Address + other info
- What happens when fully-associative is too slow?
 - Put a small (4-16 entry) direct-mapped cache in front
 - Called a “TLB Slice”
- Example for MIPS R3000:

Virtual Address	Physical Address	Dirty	Ref	Valid	Access	ASID
0xFA00	0x0003	Y	N	Y	R/W	34
0x0040	0x0010	N	Y	Y	R	0
0x0041	0x0011	N	Y	Y	R	0

Example: R3000 Pipeline Includes TLB Stages


MIPS R3000 Pipeline



TLB

64 entry, on-chip, fully associative, software TLB fault handler

Virtual Address Space



0xx User segment (caching based on PT/TLB entry)
100 Kernel physical space, cached
101 Kernel physical space, uncached
11x Kernel virtual space

Allows context switching among
64 user processes without TLB flush

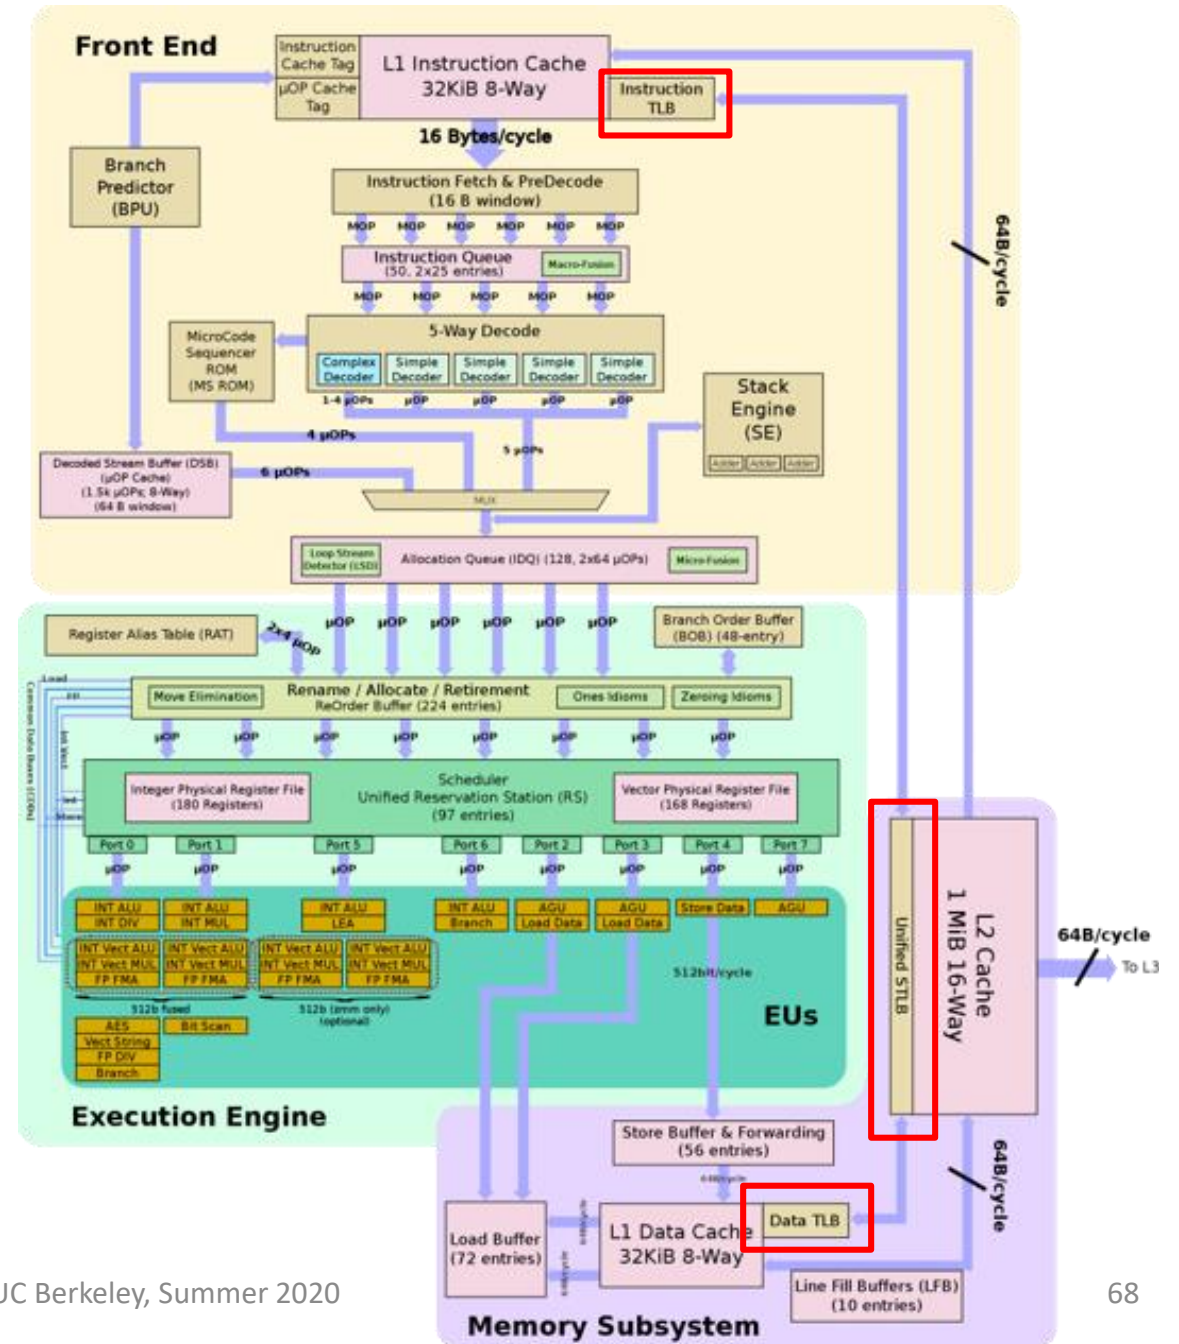
Example: Pentium-M TLBs (2003)

- Four different TLBs
 - Instruction TLB for 4K pages
 - 128 entries, 4-way set associative
 - Instruction TLB for large pages
 - 2 entries, fully associative
 - Data TLB for 4K pages
 - 128 entries, 4-way set associative
 - Data TLB for large pages
 - 8 entries, 4-way set associative
- All TLBs use LRU replacement policy
- Why different TLBs for instruction, data, and page sizes?

Example: Intel Nehalem (2008)

- L1 DTLB
 - 64 entries for 4 K pages and
 - 32 entries for 2/4 M pages,
- L1 ITLB
 - 128 entries for 4 K pages using 4-way associativity and
 - 14 fully associative entries for 2/4 MiB pages
- unified 512-entry L2 TLB for 4 KiB pages, 4-way associative

Example: Skylake, Cascade Lake



Current Example: Memory Hierarchy

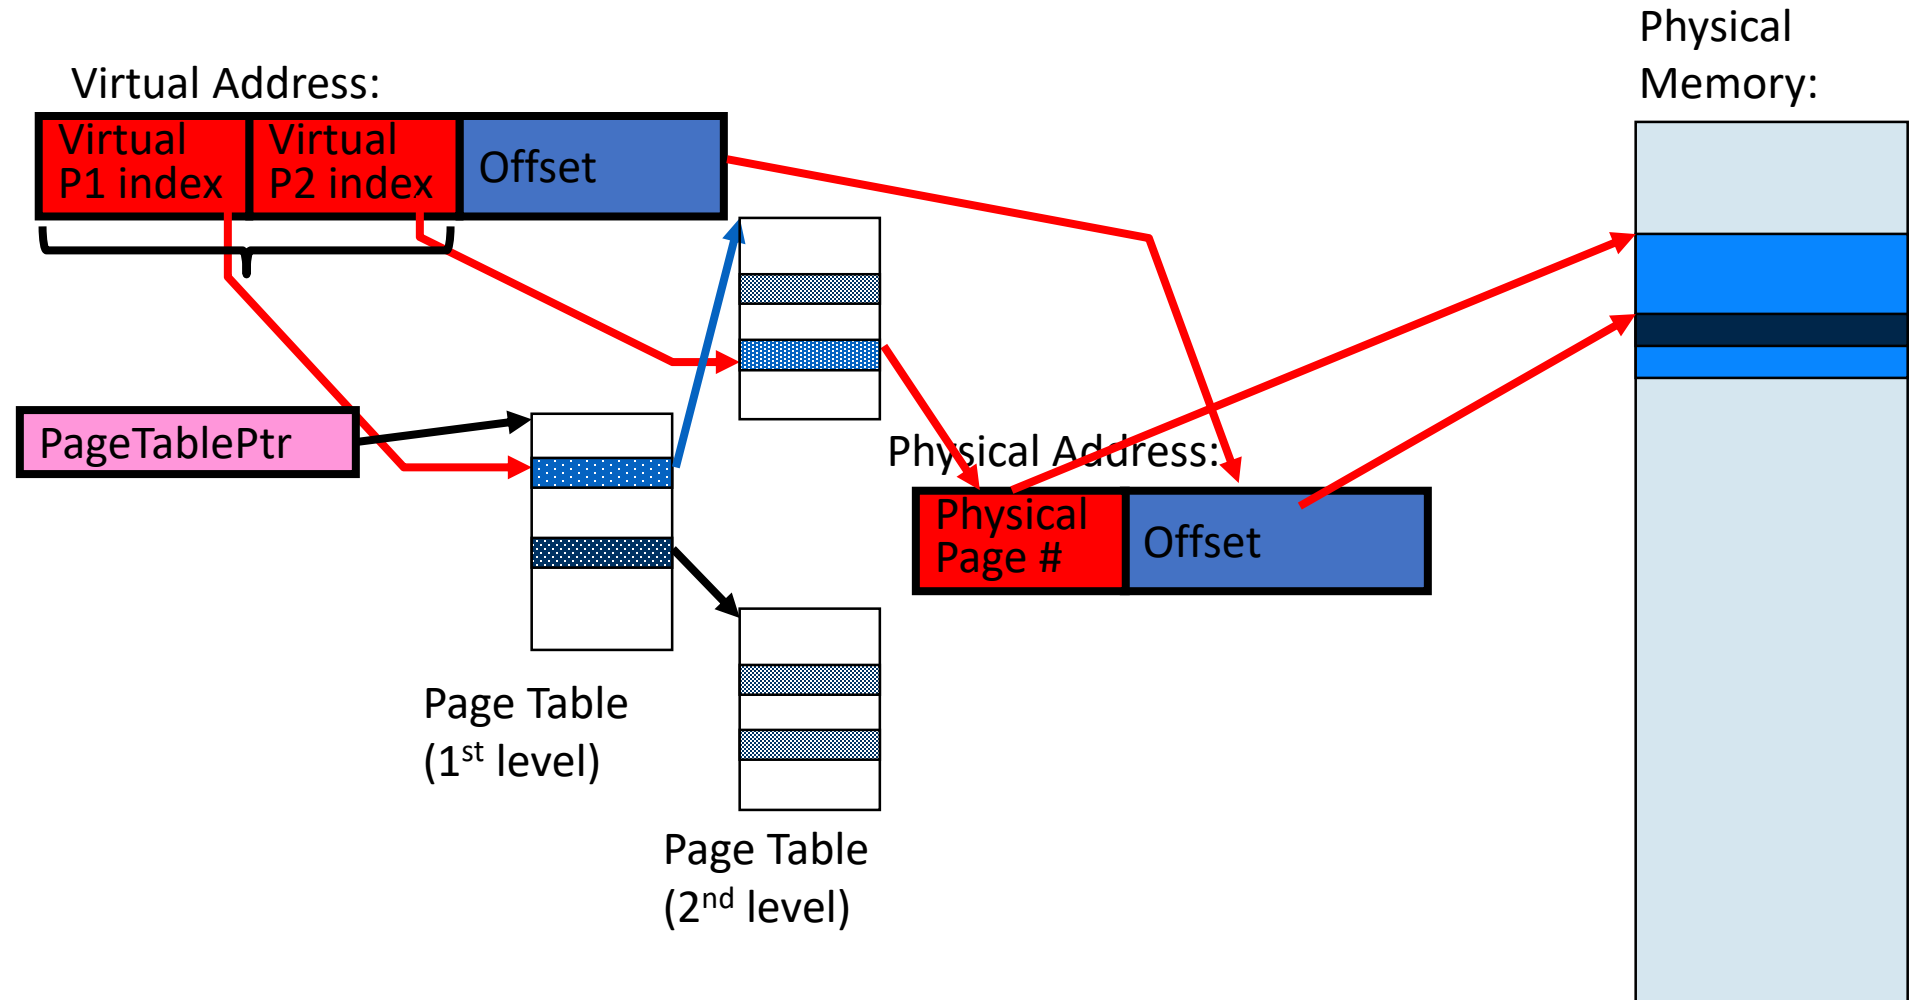
- Caches (all 64 B line size)
 - L1 I-Cache: 32 [KiB](#)/core, 8-way set assoc.
 - L1 D Cache: 32 KiB/core, 8-way set assoc., 4-5 cycles load-to-use, Write-back policy
 - L2 Cache: 1 MiB/core, 16-way set assoc., Inclusive, Write-back policy, 14 cycles latency
 - L3 Cache: 1.375 MiB/core, 11-way set assoc., shared across cores, Non-inclusive victim cache, Write-back policy, 50-70 cycles latency
- TLB
 - L1 ITLB, 128 entries; 8-way set assoc. for 4 KB pages
 - 8 entries per thread; fully associative, for 2 MiB / 4 MiB page
 - L1 DTLB 64 entries; 4-way set associative for 4 KB pages
 - 32 entries; 4-way set associative, 2 MiB / 4 MiB page translations
 - 4 entries; 4-way associative, 1G page translations
 - L2 STLB: 1536 entries; 12-way set assoc. 4 KiB + 2 MiB pages
 - 16 entries; 4-way set associative, 1 GiB page translations

What Happens on a Context Switch?

- Address Space just changed, so TLB entries no longer valid!
- Options?
 - Invalidate TLB (simple but expensive)
 - Include ASID (address space identifier) in TLB
- What if the OS changes the page table?
 - Must invalidate TLB entry!
 - Called “TLB Consistency”

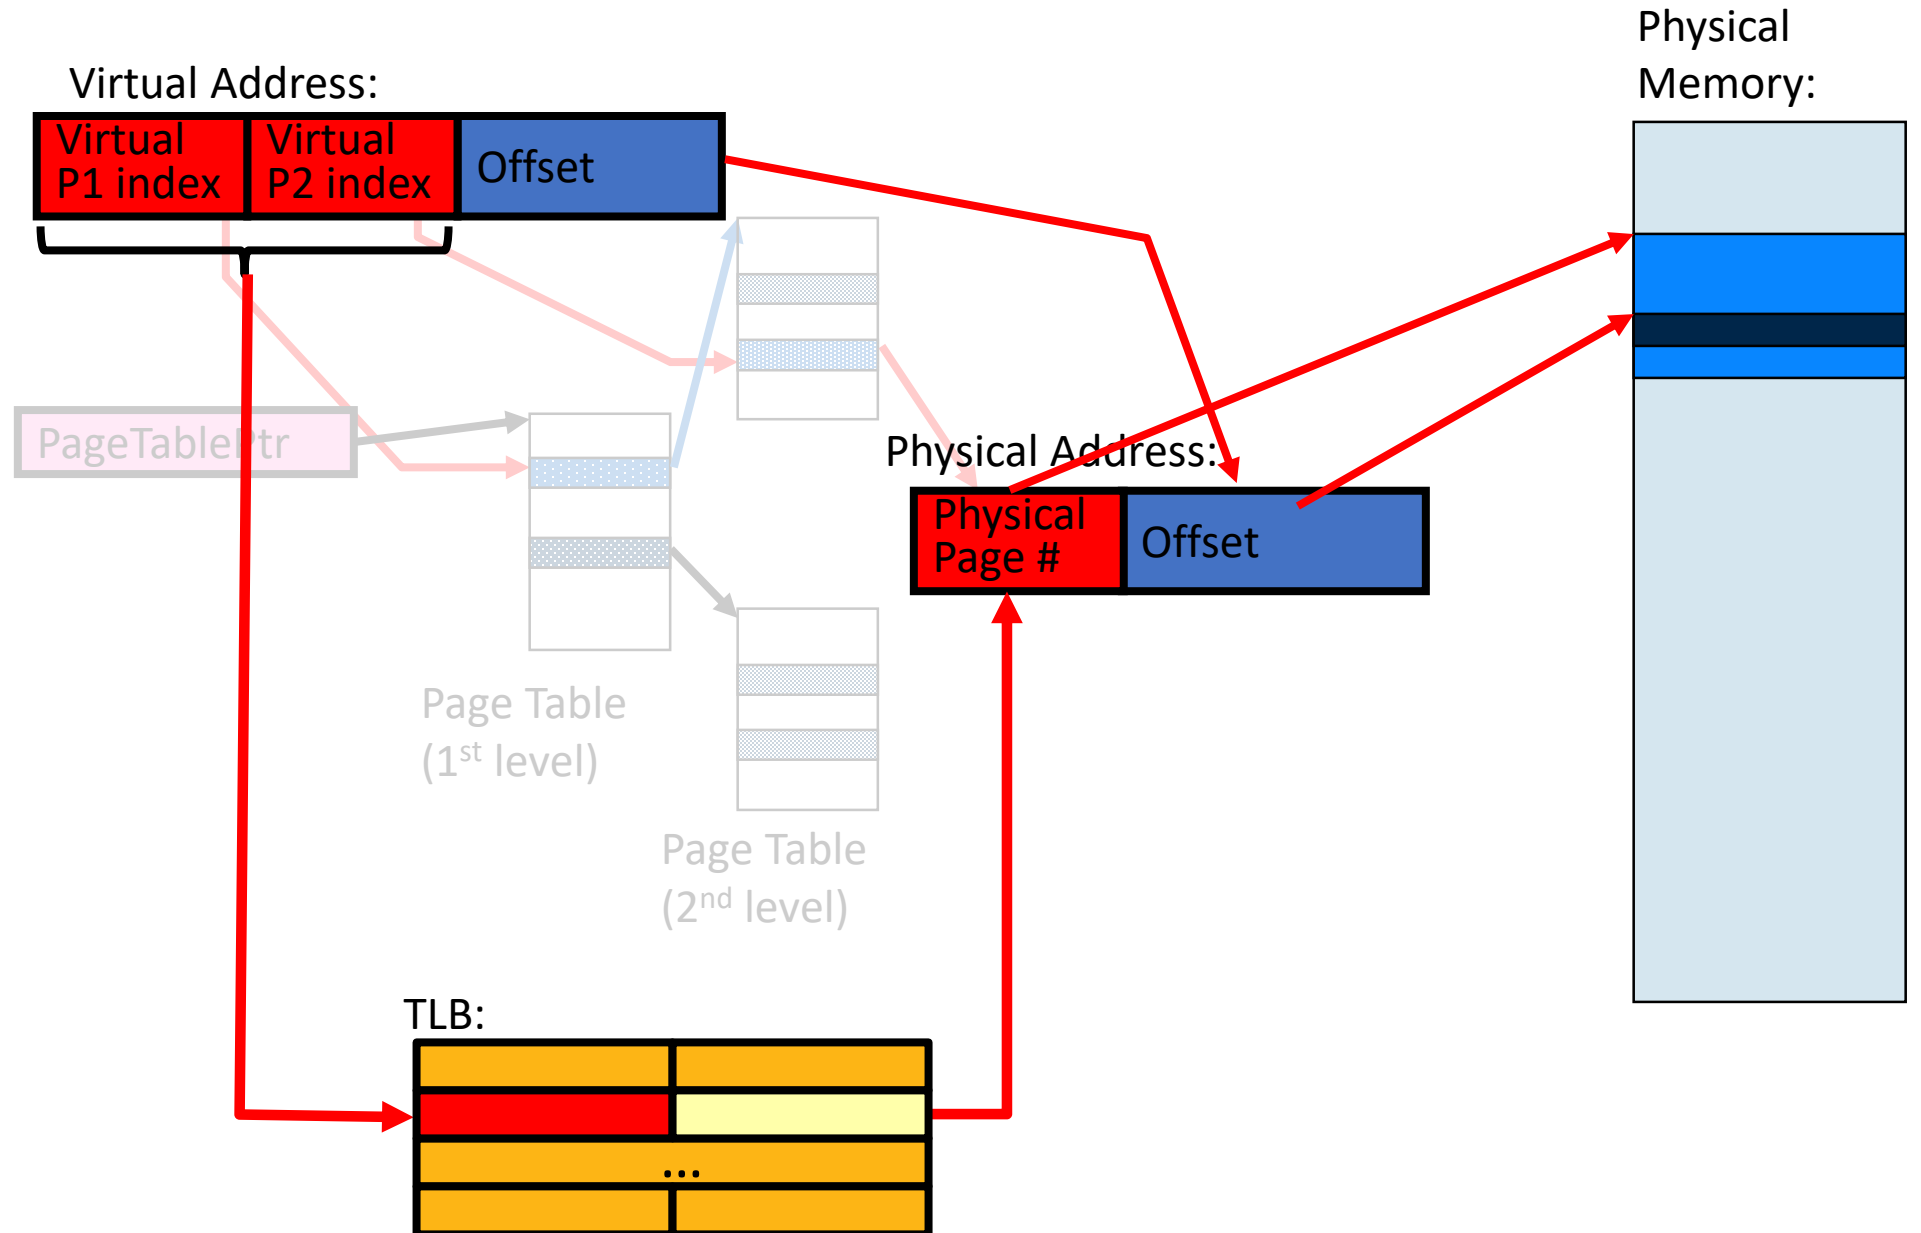
Putting Everything Together

Address Translation



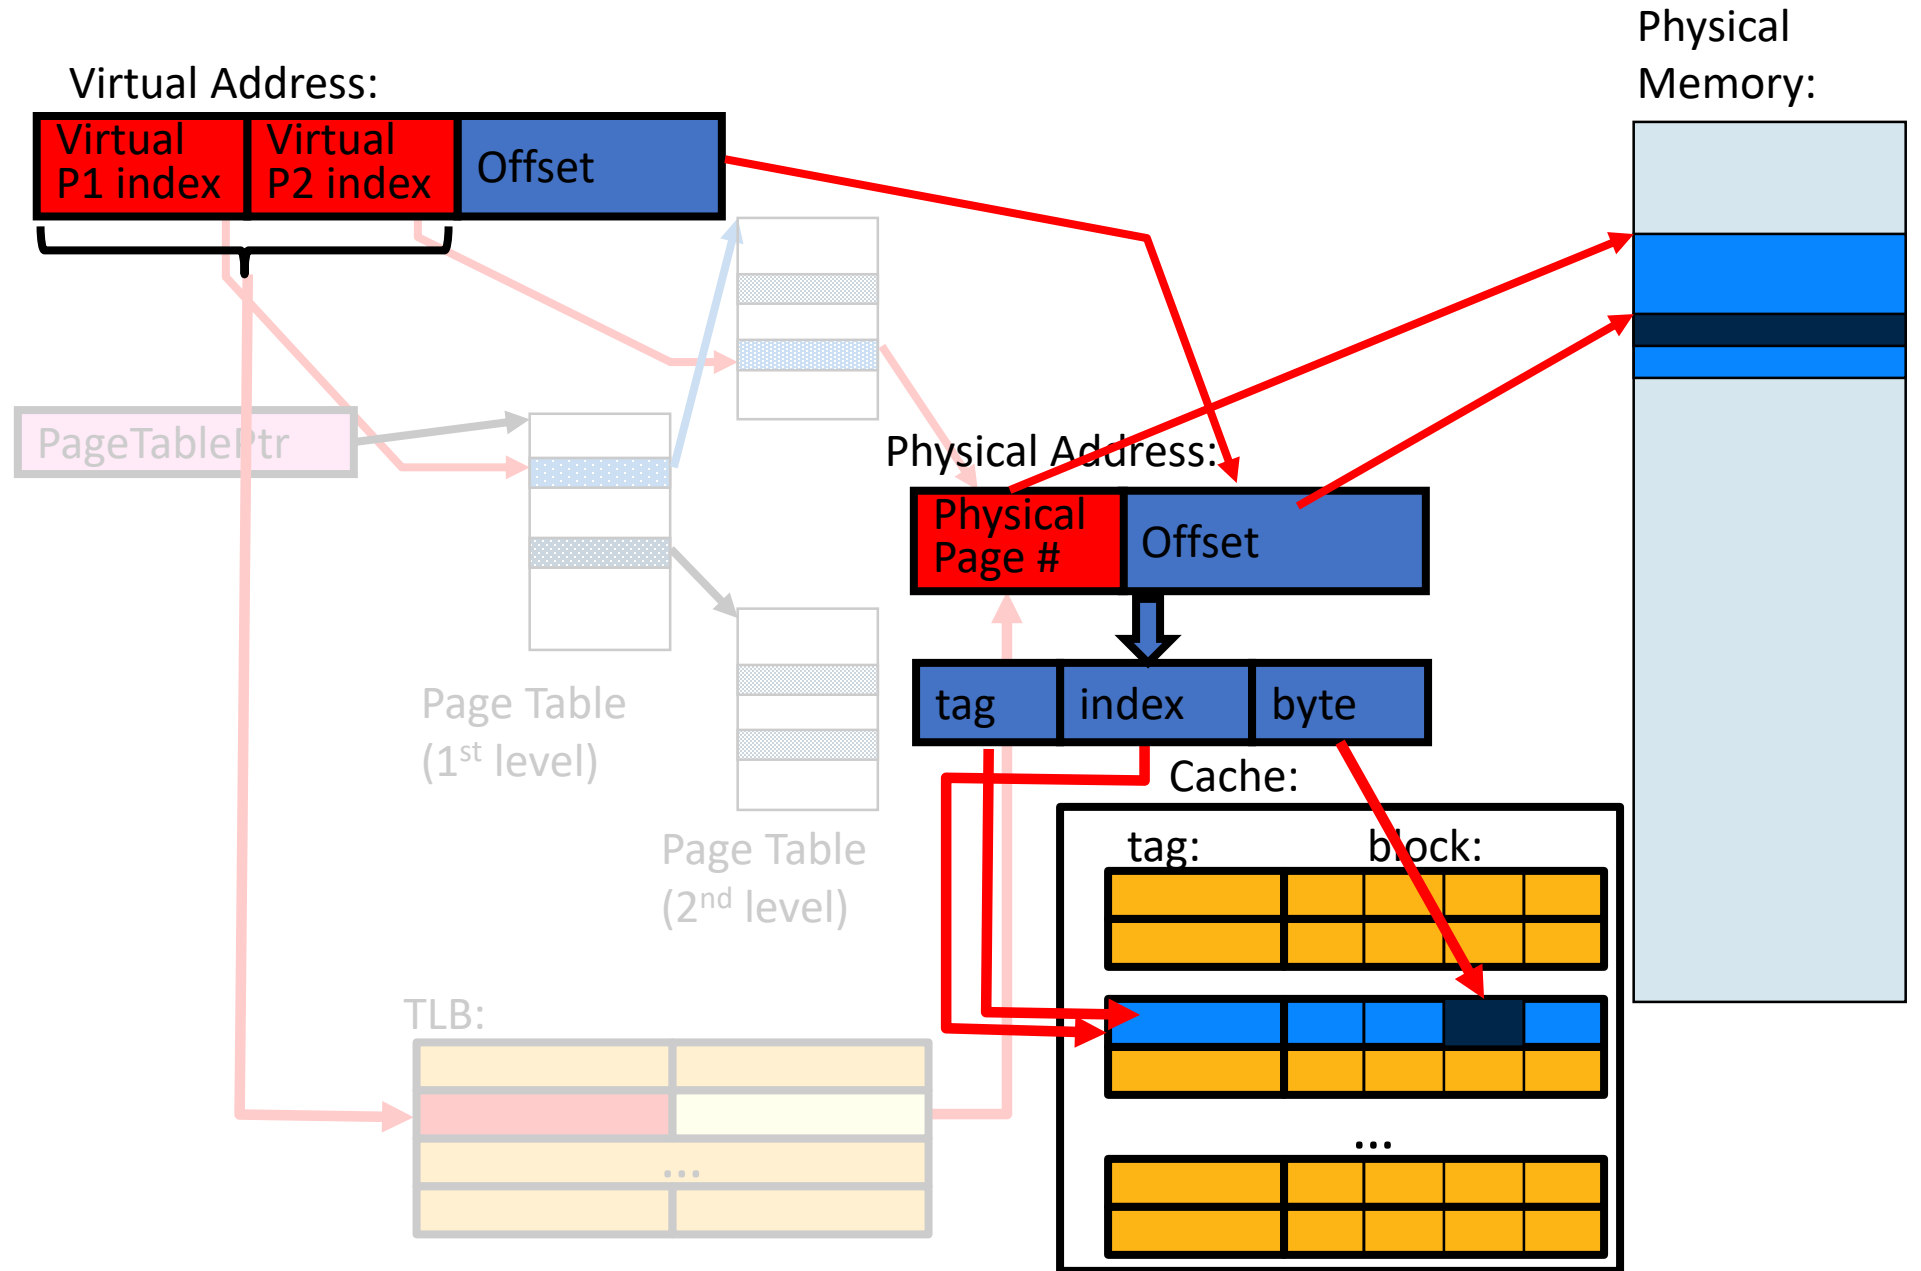
Putting Everything Together

TLB

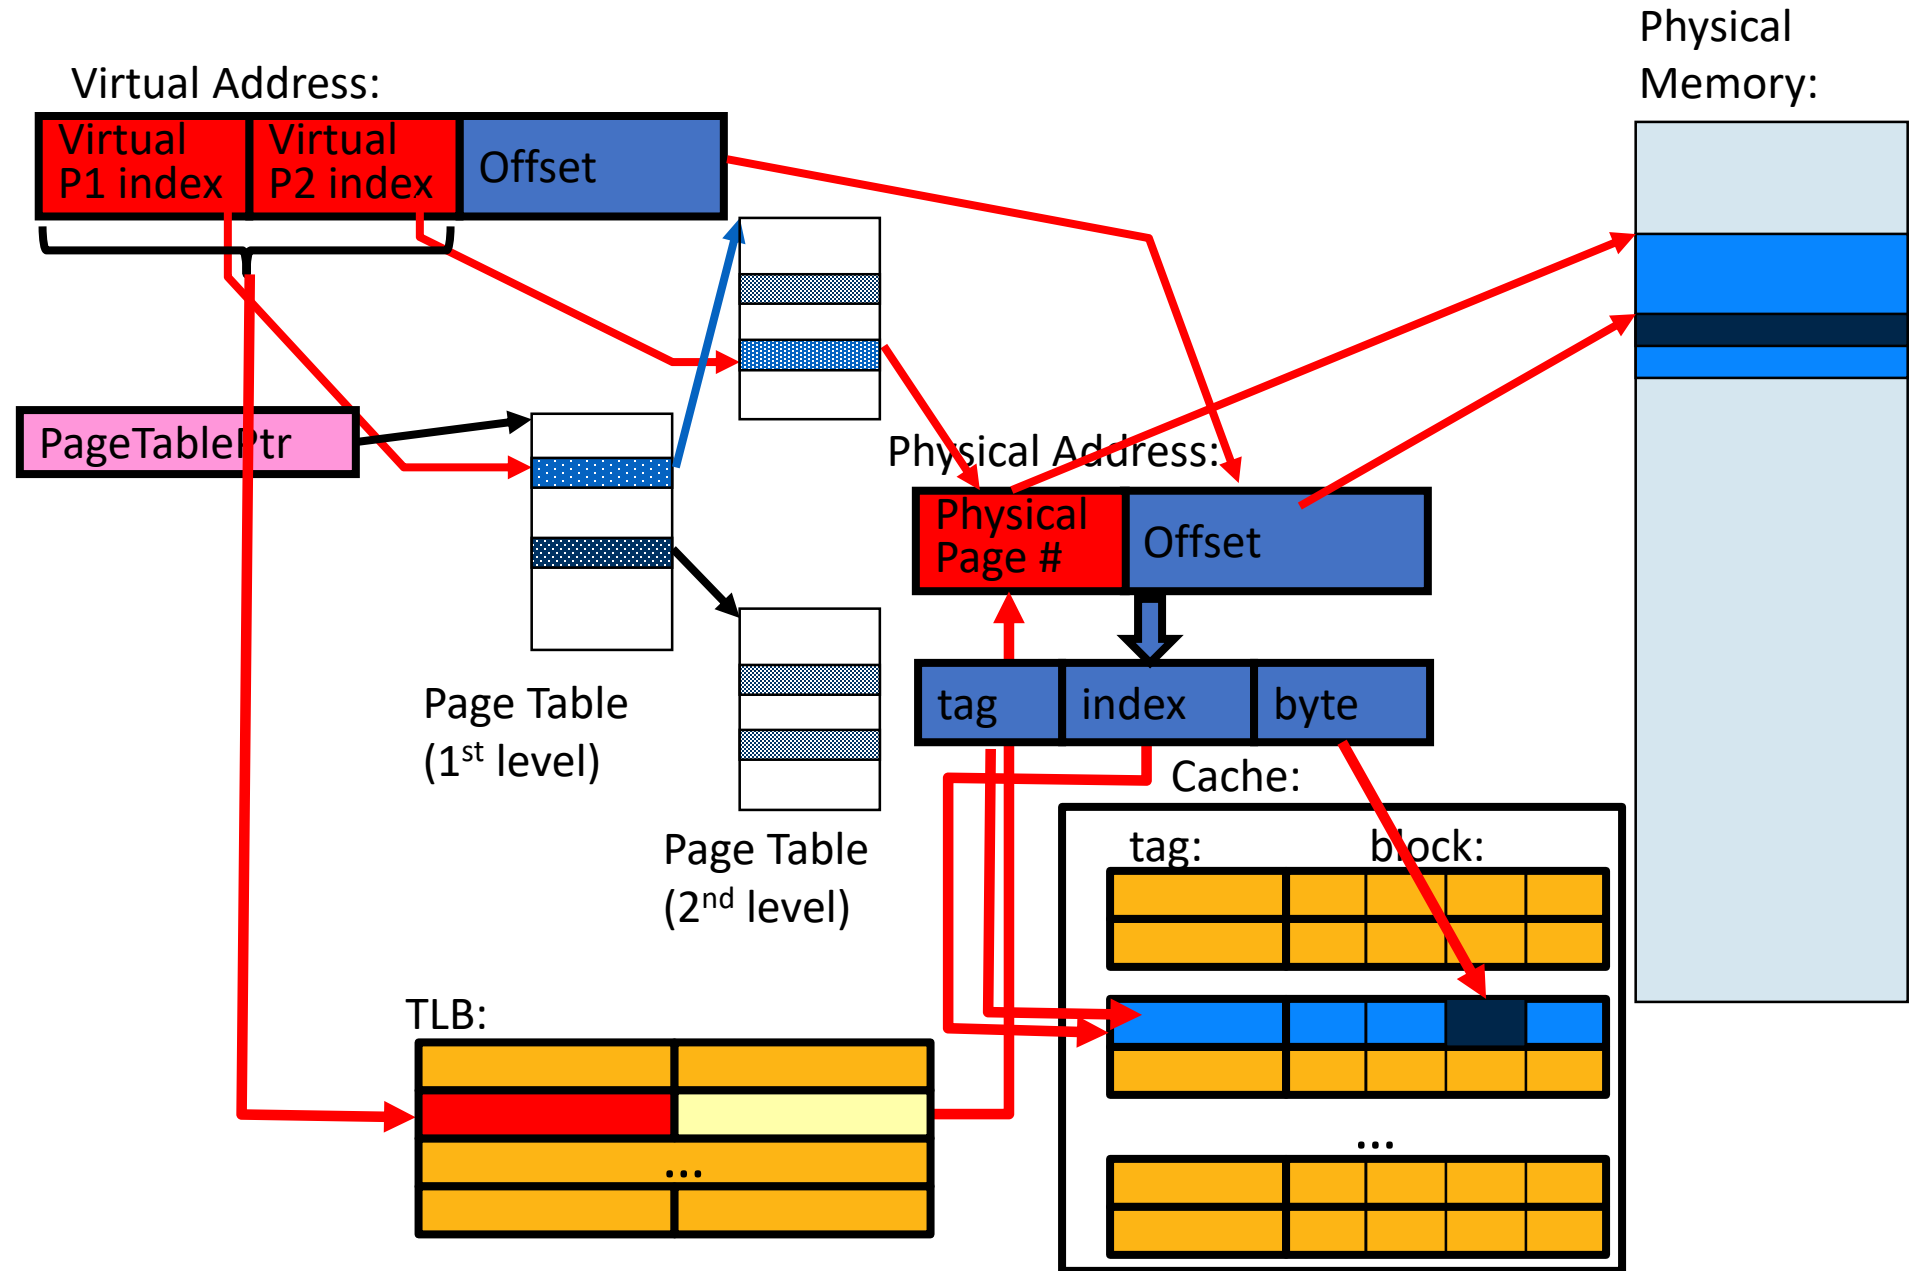


Putting Everything Together

Cache



Putting Everything Together



Summary: Page Tables

- Page Tables
 - Memory divided into fixed-sized chunks of memory
 - Virtual page number from virtual address mapped through page table to physical page number
 - Offset of virtual address same as physical address
 - Large page tables can be placed into virtual memory
- Multi-Level Tables
 - Virtual address mapped to series of tables
 - Permit sparse population of address space
- Inverted Page Table
 - Use of hash-table to hold translation entries
 - Size of page table \sim size of physical memory rather than size of virtual memory

Summary: Caching

- The Principle of Locality:
 - Program likely to access a relatively small portion of the address space at any instant of time.
 - **Temporal Locality:** Locality in Time
 - **Spatial Locality:** Locality in Space
- Three (+1) Major Categories of Cache Misses:
 - **Compulsory Misses:** sad facts of life. Example: cold start misses.
 - **Conflict Misses:** increase cache size and/or associativity
 - **Capacity Misses:** increase cache size
 - **Coherence Misses:** Caused by external processors or I/O devices
- Cache Organizations:
 - Direct Mapped: single block per set
 - Set associative: more than one block per set
 - Fully associative: all entries equivalent

Summary: TLBs

- “Translation Lookaside Buffer” (TLB)
 - Small number of PTEs and optional process IDs (< 512)
 - Fully Associative (Since conflict misses expensive)
 - On TLB miss, page table must be traversed and if located PTE is invalid, cause Page Fault
 - On change in page table, TLB entries must be invalidated
 - TLB is logically in front of cache (need to overlap with cache access)
- On Page Fault, OS can take actions to resolve the situation
 - Demand paging, automatic memory management
 - Make copy of existing page for process
 - On process start, don't have to load much of executable into memory
 - Rarely used code and data may never get paged in
- Need to handle the exception carefully